Changing UFS gear and lanes at runtime on MT6991

Why the debugfs write failed with “exit h8 state fail”, and why one line fixed it. Explained from scratch.

Written 31 October 2026, updated 1 November 2026 · MediaTek MT6991 SoC evaluation board · Android GKI kernel 6.6 (`android15-6.6`) with MediaTek vendor modules

THE WHOLE STORY IN SIX LINES

1. The board's kernel has a debugfs file, `power_mode`. Writing six digits to it (for example `332211`) should change the speed (gear), width (lanes) and mode of the link between the processor and its UFS storage chip.
2. The write woke up the UFS *controller* inside the processor, but not the UFS *chip*. The link between them stayed in **hibern8**, its deep-sleep state.
3. Before any speed change, MediaTek's driver checks that the link is awake. It waited 100 ms, the link never woke, and it printed `exit h8 state fail`.
4. The fix was one call: `ufs_debugfs_get_user_access()`, which runs `ufshcd_rpm_get_sync()`, instead of `pm_runtime_get_sync(hba->dev)`. The new call wakes the chip side too.
5. The old line was correct in Linux 4.19, where the ported code came from. In Linux 5.14 (2021) the job of waking the link moved to a different object, so the same line now does less.
6. Result on the board, 31 October: no more failures, and lane changes work.

How to read this

Every idea starts with the physical thing (wires, chips, registers) before it gets a name. You'll see six kinds of boxes, plus two markers for what was added on 1 November:

PLAIN WORDS	The simplest version of the idea.
TWO MEANINGS	A word that means different things in different places. These caused real confusion.
KEY POINT	The thing to remember from that section.
WATCH OUT	A trap you can fall into on the board.
CORRECTION	Something stated during the investigation and later corrected.
TRY IT	Commands to run on the board.
ADDED 1 NOV	New since the first version (31 October). The “What's new” list below links to every one.
UPDATE 1 NOV	Corrects or sharpens the text just above it. The original text is kept, unchanged.

Line numbers come from the copies read for this write-up: MediaTek's vendor code from a public MT6991 kernel release (`kernel_device_modules-6.6`), the Android common kernel `android15-6.6`, Qualcomm's msm-4.19 `ufs-debugfs.c`, and upstream Linux v5.13. The board's vendor tree has the same code but may be a few lines off.

ADDED 1 NOV: those copies are now in the `sources/` folder next to this file, so every line number here can be looked up (Part 11).

Contents

1. The hardware: two chips and some wires 3
2. Sleep: what hibern8 is 9
3. The software: who's who in Linux 12
4. Why the original line was right once 20
5. What happened on the board 22
6. The fix, and the other changes 27
7. How we got here, and where we went wrong
8. How we know it's fixed 3735
9. Open items 40
10. Glossary 41
11. Files in this folder 44
12. References 45

What's new on 1 November

After the first version, follow-up questions kept coming, and the answers are now in the document. Nothing was removed. For readers of the first version, these are the parts to read:

1.1 The chip's two power supplies, and who switches them 4
1.3 The clock rule only works one way 5
1.6 The PMC in the 6.6 code, call by call, and its three failure messages 7
2.2 Why writing 0 to AHIT can't wake software hibern8, and why there are two kinds 9
3.1 Which side calls the suspend and resume hooks 12
3.3 “Suspend” and “resume” are Linux words 14
3.4 What MediaTek does in its two suspend hooks 15
3.4 The controller going to sleep, and MediaTek's wrapper 18
3.5 Update: two details in the sleep picture 18
3.6 The power ladder on the board (new section) 19
5.2 Why the PMC timed out, instead of failing at once 22
5.3 Why stopping the timer couldn't help 25
5.4 The bug, row by row 25
5.5 Update, and why normal disk use never hit this bug 26
6.1 The rule the fix follows, and who else follows it 28
6.2 Is the whole function OK? A review 29
6.3 Update: the clock block, looked at again 31
6.4 Update: echo 4 is a lock, and what “only by user” means 33
7.2 More corrections 36
8.1 Reported since: fio 37
8.5 Test 3: without echo 4 (new section) 38
9 New open items 40
10 More glossary terms 42
11 New files: sources/ and the first version 44
12 New references 45

1 The hardware: two chips and some wires

1.1 What is physically on the board

Two chips matter here.

- **The MT6991 processor (the SoC).** Besides the CPU cores, it contains a block of logic called the **UFS host controller**. In this document it's **the controller**.
- **The UFS storage chip.** NAND flash memory plus its own small processor that manages the flash. In this document it's **the chip**.

Copper traces on the circuit board connect the two. These traces are the **lanes**. A lane carries bits in one direction only. Each lane is a *pair* of wires, and a bit is sent as the voltage *difference* between them: one wire goes up while the other goes down. Noise that hits both wires equally cancels out, which is why this works at billions of bits per second. MT6991 uses two lanes from the controller to the chip (transmit, TX) and two lanes back (receive, RX). That's eight data wires. The standard that defines these wires and their voltages is called **M-PHY**.

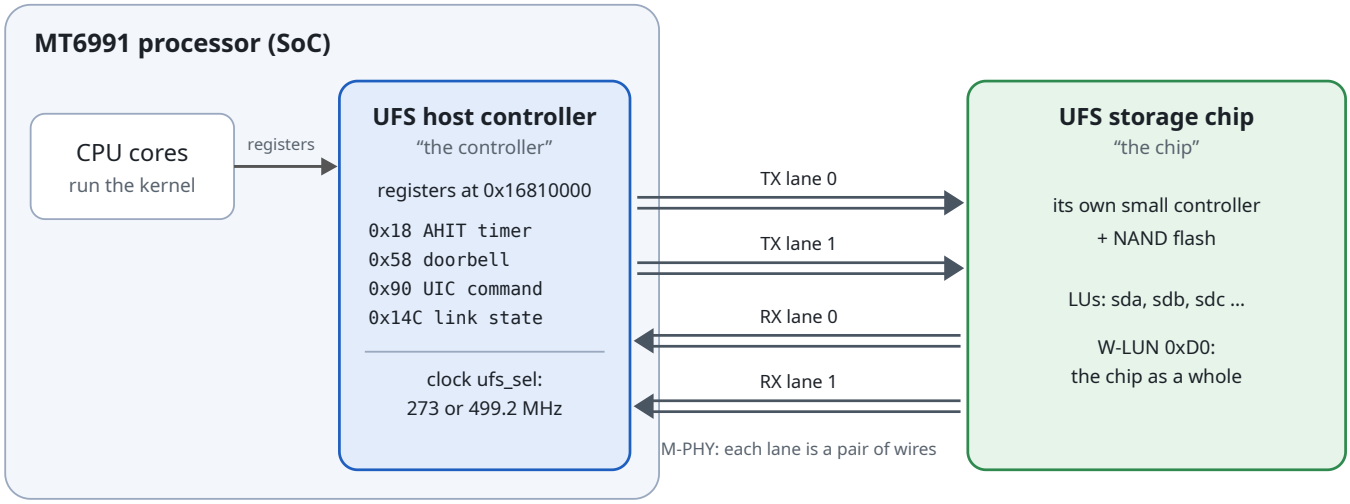


Figure 1. The controller is inside the SoC; the chip is a separate package. The CPU never touches the lanes. It only reads and writes the controller's registers, and the controller drives the wires.

The CPU talks to the controller through **registers**. A register is a memory address wired to hardware instead of to RAM. Writing a number there makes the hardware do something; reading it tells you what the hardware is doing. The controller's registers start at physical address `0x16810000`, the `ufshci@16810000` node in the device tree (`mt6991.dts:6686`). Four of them matter in this story (offsets from that base):

OFFSET	NAME	WHAT IT DOES
`0x18`	AHIT (auto-hibernate idle timer)	How long the link may sit idle before the controller puts it to sleep by itself. Writing 0 turns this off. (Part 2.)
`0x58`	Transfer request doorbell	One bit per command slot. The driver sets a bit to say "slot *n* has a command, go"; the controller clears it when the command is done.

OFFSET	NAME	WHAT IT DOES
`0x90`	UIC command	Where the driver writes commands for the link itself: wake it, put it to sleep, read or change a link setting.
`0x14C`	MediaTek link state (`REG_UFS_UFS_MMIO_OTSD_CTRL`)	MediaTek-specific. Its top byte says whether the link is up, down or asleep. (Part 2.)

TWO MEANINGS

"Host" can mean the controller hardware, Linux's SCSI host object (`struct Scsi_Host`), or loosely the whole phone. **"Device"** can mean the UFS chip, or any `struct device` in Linux's tree of objects. In this document, *the controller* and *the chip* are always the hardware. A plain "device" means the Linux object.

ADDED 1 NOV · THE CHIP'S TWO POWER SUPPLIES

The chip runs on two supply rails. They come from the **PMIC**, a separate power-management chip on the board (an MT6363 here), which the processor commands:

RAIL	FEEDS	ON THE BOARD
VCC	The flash memory itself	`vcc-supply = <&mt6363_vemc>` (`mt6991.dts:6752`)
VCCQ (1.2 V)	The chip's own small controller, and its end of the link	`vccq-supply = <&mt6363_vufs12>` (`mt6991.dts:6753`)

Only the host switches these rails; the chip has no say. All the chip gets is a warning first (START STOP UNIT, Part 3). Section 3.6 shows the moment each rail changes.

1.2 Three layers of rules on the wires

Wires alone aren't enough; both ends need rules. UFS stacks three sets of them, and the chip runs the same stack on its side:

UFS transport (UTP)	Carries SCSI commands (READ, WRITE, START STOP UNIT, …) in packets called UPIUs. This is what the files are made of.
UniPro	The link layer. Packs bits into frames, checks them, acknowledges them and re-sends anything lost. It also owns the link's *speed and power state*. Its settings are numbered **attributes**, such as `PA_TXGEAR` (0x1568, the transmit gear) or `PA_ACTIVETXDATALANES` (0x1560, how many transmit lanes are in use). "PA" stands for PHY Adapter, the bottom part of UniPro.
M-PHY	The electrical layer: drives the voltages, recovers bits and timing at the other end, puts the wires to sleep and wakes them.

The driver reaches UniPro through **UIC commands**, written to register 0x90. The ones in this story are `DME_GET` and `DME_SET` (read or write one attribute), and `DME_HIBERNATE_ENTER` and `DME_HIBERNATE_EXIT` (put the link to sleep, wake it up). DME is the part of UniPro that accepts these commands (the Device Management Entity).

> **PLAIN WORDS**
>
> SCSI is what the letter says. UniPro is the delivery service that numbers the letters, confirms delivery and re-sends lost ones. M-PHY is the road.

1.3 Gears: how fast the wires flip

A **gear** is how many bits per second each lane carries. The high-speed (HS) gears go from G1 to G5, and each gear doubles the one before. There are two slightly different base frequencies, called **rate A** and **rate B**. MediaTek always uses rate B (it's the default, `PA_HS_MODE_B`, in the core's shared negotiation helper).

GEAR	BITS PER SECOND PER LANE (RATE B)	ONE BIT LASTS
HS-G1	1,457.6 million	686 ps
HS-G2	2,915.2 million	343 ps
HS-G3	5,830.4 million	172 ps
HS-G4	11,660.8 million	86 ps
HS-G5	23,321.6 million	43 ps

To make that physical: at G5 one bit lasts 43 picoseconds (trillionths of a second). In that time a signal travelling along a circuit-board trace covers about 6 or 7 millimetres.

Not every bit on the wire is your data. Gears G1 to G4 send every 8 bits of data as 10 bits on the wire (**8b/10b** encoding), which keeps the signal balanced and lets the receiver stay locked to its timing. G5 uses the more efficient 128b/130b. So the data a lane can carry is its bit rate × 0.8 (× 128/130 at G5). Part 8 uses this to check that a gear change really happened.

There is also a slow family, the **PWM gears**, used for bring-up and low power. They run from a few megabits per second to a few hundred.

The controller's own clock has to keep up with the gear. On MT6991 that clock is `ufs_sel`, which switches between a 273 MHz source and a 499.2 MHz source (`freq-table-hz = <273000000 499200000>`, `mt6991.dts:6732`). MediaTek's code pairs the top gear with 499.2 MHz and G4 with 273 MHz (section 6.4). The 499.2 MHz source needs the SoC's core supply voltage (vcore) at 0.65 V or more (`clk-scale-up-vcore-min = <650000>`, `mt6991.dts:6751`).

> **ADDED 1 NOV · THE CLOCK RULE ONLY WORKS ONE WAY**
>
> "Has to keep up" means the controller's clock must never be *slower* than the gear needs. Faster is always fine; slowing it down at low gears only saves power.
>
> - **The gear alone sets the wire speed.** A PMC sets `PA_TXGEAR`, and the M-PHY switches its line rate. The link has its own clocks for that, listed in the device tree as `unipro_tx_sym`, `unipro_rx_sym0` and

`unipro_rx_sym1` (`mt6991.dts:6723`). There is one receive clock per lane because each lane's receiver recovers its timing from the bits arriving on it.

- **`ufs_sel` is the controller's clock** (`mt6991.dts:6717`): it sets how fast the controller can move data between memory and the link. A fast controller in front of a slow link just waits more.
- **Every boot already does this.** The core sets `ufs_sel` to 499.2 MHz in `ufshcd_init_clocks()` (`ufshcd.c:9398`) before the link starts, and the link then runs its first commands in PWM-G1, the slowest mode there is.
- **G1 to G3.** MediaTek's code never chooses G2 or G3; the only lower gear it picks on its own is G4, on the 273 MHz source. Each gear doubles the one before, so G1 to G3 carry less than G4. 273 MHz is enough for them, and 499.2 MHz is more than enough for every gear.
- **vcore follows the clock, not the gear.** 499.2 MHz needs at least 0.65 V at any gear; 273 MHz doesn't. MediaTek's helper `_ufs_mtk_clk_scale()` (`ufs-mediatek.c:3039`) always moves the two together: vcore up before the clock goes up, and down only after the clock comes down.

The one place MediaTek itself runs HS-G1 is its eye-monitor and scrambler debug code (`ufs-mediatek-sysfs.c:307` and `614`), and there it uses the 273 MHz source. Section 6.3 uses this rule.

1.4 Lanes

Lanes are parallel roads. With two lanes the data is split across both, so the link carries twice as much. Each direction can use 1 or 2 lanes.

> **TWO MEANINGS**
>
> "Two lanes" in UFS means two *per direction*: 2 TX and 2 RX, four wire pairs. In the six digits, the lane part `22` is one digit for RX and one for TX.

1.5 Modes: the last two digits

- `1` = FAST_MODE: the HS gears.
- `2` = SLOW_MODE: the PWM gears.
- `4` = FASTAUTO, `5` = SLOWAUTO: the same, except that the link may drop into a lower-power state on its own between bursts of data.

> **TWO MEANINGS: "POWER MODE" HAS THREE**
>
> 1. **Link power mode** (UniPro): FAST, SLOW, FASTAUTO, SLOWAUTO. The link's speed setting. The digits set this.
> 2. **Chip power mode**: Active, Sleep, PowerDown. Set by sending the chip the SCSI command START STOP UNIT (often shortened to SSU).
> 3. **Linux runtime PM state** of a `struct device`: active or suspended (Part 3).
>
> The bug involved all three at once.

1.6 A power mode change, step by step

Changing gear, lanes or mode is a **power mode change**, or **PMC**. Both ends must switch at the same moment, so a PMC is a conversation over the link:

1. The driver writes the new settings into the controller's UniPro attributes: `PA_RXGEAR`, `PA_ACTIVERXDATALANES`, `PA_RXTERMINATION`, `PA_TXGEAR`, `PA_ACTIVETXDATALANES`, `PA_TXTERMINATION`, `PA_HSSERIES` (rate A or B), and some timeout values (`PA_PWRMODEUSERDATA0` and onwards).
2. It writes `PA_PWRMODE`. That write is the trigger.
3. UniPro negotiates the change with the chip, over the link.
4. The controller raises an interrupt with the result. “Local success” means done.

In the 6.6 code, `ufshcd_change_power_mode()` (`ufshcd.c:4560`) does steps 1 and 2, and `ufshcd_uic_pwr_ctrl()` (`ufshcd.c:4199`) waits up to 500 ms (`UIC_CMD_TIMEOUT_DEFAULT`) for step 4. If nothing arrives, it prints `pwr ctrl cmd 0x… with mode 0x… completion timeout` (`ufshcd.c:4237`).

ADDED 1 NOV · THE PMC IN THE 6.6 CODE, CALL BY CALL

The core file is now in this folder at `sources/android15-6.6/ufshcd.c`, so you can follow along. Line numbers are on the right:

```
ufshcd_config_pwr_mode()                 4652  the ported function calls this
 ├─ vendor hook, PRE_CHANGE              ufshcd-priv.h:182
 │    → MediaTek ufs_mtk_pwr_change_notify()  ufs-mediatek.c:2023
 │      saves AHIT, turns auto-hibern8 off, picks the final gear/lanes/mode
 └─ ufshcd_change_power_mode()           4560
     ├─ "power already configured": return, nothing sent   4566
     ├─ DME_SETs: PA_RXGEAR, PA_TXGEAR, lanes, termination, HS series,
     │   timeouts. Each goes ufshcd_dme_set_attr() 4074 → __ufshcd_send_uic_cmd() 2468
     ├─ ufshcd_uic_change_pwr_mode()     4333  writes PA_PWRMODE: the actual "go"
     │   └─ ufshcd_uic_pwr_ctrl()        4199  sends it, then sleeps on uic_async_done
     │        for up to 500 ms (UIC_CMD_TIMEOUT_DEFAULT, line 59).
     │        Woken by the interrupt handler ufshcd_uic_cmd_compl() 5472,
     │        which calls complete() at 5493 when the power-mode interrupt arrives.
     └─ only on success: vendor hook POST_CHANGE (MediaTek writes AHIT back),
        then hba->pwr_info = the new mode   4638
```

The conversation in step 3. Writing `PA_PWRMODE` makes the host's UniPro send a power-change request across the wires and wait for the chip's confirmation. UniPro calls these frames `PACP_PWR_req` and `PACP_PWR_cnf`. When that exchange finishes, whether it worked or not, the controller raises the interrupt that `ufshcd_uic_pwr_ctrl()` is waiting for (`ufshcd.c:4234`).

It can fail in three ways, and each prints something different:

MESSAGE	LINE	WHAT IT MEANS
`pwr ctrl cmd … uic error`	4229	The host's own UniPro refused the command. You see it at once.
`pwr ctrl cmd … failed, host upmcrs:…`	4254	The exchange finished, but the result wasn't “local success”.

MESSAGE	LINE	WHAT IT MEANS
`pwr ctrl cmd … completion timeout`	4237	The exchange never finished. The host accepted the command and heard nothing back for 500 ms. This was Error 1 (5.2).

After any of them, unless a suspend or resume is in progress, the core marks the link broken and schedules the error handler (`ufshcd.c:4270` to `4272`). That resets the controller and brings the link back at its maximum gear.

KEY POINT

A PMC is a conversation over the link. If the link is asleep, nobody answers.

The six digits are `GGLLMM`: gear RX, gear TX, lanes RX, lanes TX, mode RX, mode TX. So `332211` means HS-G3, two lanes each way, FAST mode.

2 Sleep: what hibern8 is

2.1 The link's deep sleep

When there's nothing to send, keeping eight wires switching billions of times a second wastes power. **Hibern8** (say “hibernate”) is the link's deepest sleep that still remembers its settings. The transmitters stop driving the wires, and the circuits that generate the fast clocks (the PLLs) stop. The gear and lane settings are kept, so waking up is much quicker than starting the link from scratch.

> **TWO MEANINGS**
>
> Hibern8 is not the PC kind of hibernate, where RAM is saved to disk. Only the link sleeps; the phone stays on.

2.2 Two ways in, and one way out of each

Auto-hibern8 (AH8): the hardware does it. You put an idle time into the AHIT register (offset 0x18). If no command arrives for that long, the controller puts the link to sleep by itself, and when a new command arrives it wakes the link by itself. Software takes no part. MediaTek sets the idle time according to who made the chip, between 1 and 3.5 ms (`ufs_mtk_fix_ahit()`, `ufs-mediatek.c:1488`). Writing 0 to AHIT turns the timer off.

Software hibern8: the driver does it. The driver sends the UIC command `DME_HIBERNATE_ENTER`. The link then stays asleep until the driver sends `DME_HIBERNATE_EXIT`. The hardware will not wake it on its own, and changing AHIT doesn't touch it.

> **PLAIN WORDS**
>
> Auto-hibern8 is a motion-sensor light: it goes off when nobody's around and comes back on when someone walks in. Software hibern8 is a light somebody switched off by hand: it stays off until somebody switches it back on.

> **KEY POINT**
>
> The whole bug in one sentence: MediaTek's check treated a hand-switched light as if it were a motion-sensor light, and waited for it to come on by itself.

> **ADDED 1 NOV · WHY WRITING 0 TO AHIT CAN'T WAKE SOFTWARE HIBERN8**
>
> - **Auto-hibern8 is the controller's own small state machine.** The timer runs out and the controller takes the link down; a doorbell, or AHIT going to 0, makes it bring the link back. It only undoes what it did itself.
> - **Software hibern8 sits outside that state machine.** A UIC command, `DME_HIBERNATE_ENTER`, took the link down, and only the opposite command brings it back. In 6.6 that's sent by `ufshcd_uic_hibern8_exit()` (`ufshcd.c:4413`) and, on MediaTek, also by `ufs_mtk_link_set_hpm()` (`ufs-mediatek.c:2467`) in its resume hook. Both run only when the W-LUN wakes.

- **On the board AHIT is already 0 by then.** MediaTek's suspend hook writes 0 before the chip goes to Sleep (3.4), and nothing writes it back until the W-LUN wakes. Writing 0 again doesn't even change the register, so the controller has nothing to react to.
- **Both ends of the link go down, not the chip.** Hibern8 also takes down the chip's transmitters and receivers, its half of the lanes. The chip itself keeps running: its controller is on, and VCC and VCCQ stay on.

Why there are two kinds

They do different jobs on different time scales:

- **Auto-hibern8 covers the short gaps while the disk is in use.** After 1 ms of quiet the controller takes the link down, and the next doorbell brings it back, with no CPU, interrupt or driver code involved. Doing that in software would take two UIC commands for every gap, each with an interrupt round trip.
- **Software hibern8 is for when more than the link goes down:** the chip into Sleep, the M-PHY powered off, clocks off. Before cutting power, software must know the link is in hibern8 (it waited for the command to finish) and that it will stay there. Auto-hibern8 can promise neither: the link can be up or down at any moment, and the next doorbell wakes it.

They take turns and never overlap. The timer is off whenever software drives the link, and it comes back when software hands the link back:

- **Suspend and resume:** MediaTek's suspend hook turns it off (3.4), and the core writes it back as the last step of the W-LUN's wake (`ufshcd.c:9999`).
- **A PMC:** MediaTek's `pwr_change_notify` turns it off before and writes it back after (`ufs-mediatek.c:2023`).
- **Clock gating:** MediaTek lets the core enter software hibern8 during clock gating only when auto-hibern8 is disabled (`ufs-mediatek.c:1779`). On the board it isn't, so clock gating never uses software hibern8.

2.3 How the driver sees whether the link is awake

MediaTek's controller reports the link's state in register `0x14C` (`REG_UFS_UFS_MMIO_OTSD_CTRL` , `ufs-mediatek.h:52`). Its top byte, bits 31 to 24, holds the state:

VALUE	NAME IN THE CODE	MEANING
0	`VS_LINK_DISABLED`	Link not started
1	`VS_LINK_DOWN`	Link down
2	`VS_LINK_UP`	Awake and working
3	`VS_LINK_HIBERN8`	Asleep in hibern8
4	`VS_LINK_LOST`	Connection lost
5	`VS_LINK_CFG`	Being configured

`ufs_mtk_wait_link_state()` (`ufs-mediatek.c:562`) reads that byte every 100 to 200 µs until it shows the state it wants, or gives up when its time limit runs out. Bits 20 to 16 of the same register show the controller's internal state machine; MediaTek's idle wait uses them to spot a hibern8 entry or exit that's still in progress (values 8 to 13).

```
static int ufs_mtk_wait_link_state(struct ufs_hba *hba, u32 state,
                                   unsigned long max_wait_ms)
{
        ...
        timeout = ktime_add_ms(ktime_get(), max_wait_ms);
        do {
                time_checked = ktime_get();
                if (host->legacy_ip_ver || host->ip_ver < IP_VER_MT6899) { ... older chips ... }
                else {
                        val = ufshcd_readl(hba, REG_UFS_UFS_MMIO_OTSD_CTRL);
                        val = val >> 24;
                }
                if (val == state)
                        return 0;
                usleep_range(100, 200);   /* Sleep for max. 200us */
        } while (ktime_before(time_checked, timeout));

        return -ETIMEDOUT;
}
```

ufs-mediatek.c:562 (trimmed)

MT6991 uses this register because its IP version (`0x10460000` or `0x10470000`) is at least `IP_VER_MT6899`. Older MediaTek chips read `REG_UFS_PROBE` instead. An earlier explanation called it `REG_UFS_PROBE`; that's the older chips' register.

3 The software: who's who in Linux

3.1 Core and vendor: where the ported code runs

On Android, one generic kernel image is built for every phone, the **GKI** (Generic Kernel Image). The UFS **core** is part of it: `drivers/ufs/core/ufshcd.c`, `ufs-debugfs.c` and their friends. MediaTek's UFS code (`ufs-mediatek.c`, `ufs-mediatek-sysfs.c`, ...) is a **vendor module** loaded on top, from `kernel_device_modules`.

The core calls into the vendor module through **vops** ("variant operations"): a table of function pointers that MediaTek fills in. The ones in this story are:

- `pwr_change_notify`, called just before a PMC (`PRE_CHANGE`) and just after it (`POST_CHANGE`);
- `suspend` and `resume`, called when the chip side goes to sleep or wakes;
- `clk_scale_notify`, called around clock changes.

The debugfs file lives in the core's `ufs-debugfs.c`. It can call core functions, but it can't see MediaTek's private data (`struct ufs_mtk_host`). That's why the fix had to use a core helper. It's also why a change there needs a new kernel image, not just a new vendor module.

> **ADDED 1 NOV · WHICH SIDE CALLS SUSPEND AND RESUME**
>
> The hooks' names don't say whose suspend they mean. In 6.6, `suspend` and `resume` are called by the W-LUN, the chip side:
>
> - `__ufshcd_wl_suspend()` calls `suspend` twice: `PRE_CHANGE` at `ufshcd.c:9830` and `POST_CHANGE` at `9879`.
> - `__ufshcd_wl_resume()` calls `resume` at `ufshcd.c:9939`.
>
> Up to Linux 5.13 the controller's own `ufshcd_suspend()` called the same hook (`ufshcd_v5.13.c:8769`). The name stayed the same and the caller moved, together with the job (Part 4).
>
> The controller's own sleep and wake give the vendor a different hook, `setup_clocks`, called from `ufshcd_setup_clocks()` (`ufshcd.c:9298` and `9329`). It runs when the controller sleeps or wakes (`10136`, `10176`), and also on every clock gate and ungate (`1910`, `1768`). Sections 3.4 and 3.6 say what MediaTek does in each.

> **TWO MEANINGS: DEBUGFS AND SYSFS**
>
> Both are files the kernel invents on the fly. **sysfs** (`/sys`) is the stable interface for real users; MediaTek's `clkscale_control` lives there. **debugfs** (`/sys/kernel/debug`) is for developers and can change at any time; the `power_mode` file lives there.

3.2 The objects Linux creates for UFS

Linux keeps a tree of `struct device` objects, one for each thing it manages. For UFS the tree looks like this:

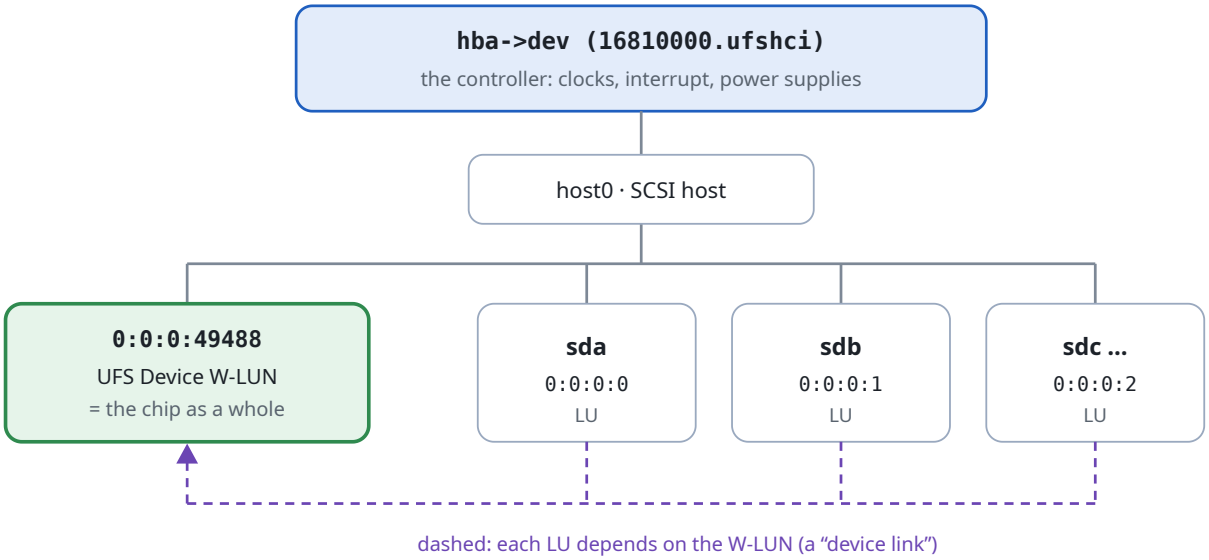


Figure 2. Solid lines are parent and child; dashed lines are “depends on”. The W-LUN is Linux's handle on the chip itself. There are two more W-LUNs (RPMB and BOOT) that don't matter here.

- `hba->dev` is the controller, as a “platform device” named `16810000.ufshci`. Its jobs: the controller's clocks, interrupt and power supplies. (`hba` stands for host bus adapter, the core's `struct ufs_hba` for one controller.)
- `host0` is Linux's SCSI host. The chip speaks SCSI, so Linux handles it like a SCSI disk enclosure.
- **Logical units (LUs)** are slices of the chip that you can read and write. They appear as `sda`, `sdb`, `sdc`, and so on; your partitions live inside them.
- **Well-known logical units (W-LUNs)** are special addresses for jobs that concern the whole chip. The important one is the **UFS Device W-LUN**: to put the whole chip to sleep or wake it, you send START STOP UNIT to this address.

The UFS Device W-LUN is number `0xD0` in UFS's own numbering. Linux shows it as `0:0:0:49488`. The arithmetic: clear the W-LUN flag bit (`0x80`), which leaves `0x50`; add SCSI's base for well-known units, `0xC100`; that gives `0xC150`, which is 49488. (The RPMB W-LUN, `0xC4`, becomes 49476, and the BOOT W-LUN, `0xB0`, becomes 49456.)

> **PLAIN WORDS**
>
> The W-LUN is Linux's handle on *the chip itself*. Waking the W-LUN means waking the chip and the link that leads to it.

3.3 Runtime PM from scratch

Runtime PM is how Linux lets single parts of a running phone sleep while everything else keeps working. Every `struct device` has a usage counter and a state, *active* or *suspended*.

- `pm_runtime_get_sync(dev)` means “I need this device now”. It adds 1 to the counter. If the device is asleep, it runs the device's wake-up function (its *resume callback*) and waits for it to finish.
- `pm_runtime_put(dev)` means “I'm done”. It subtracts 1. At zero, the device may go to sleep after its **autosuspend delay**.

Four rules decide who wakes whom. Each has a physical reason:

1. **Waking a device wakes its parent first.** You can't reach the chip if the controller's clocks are off.
2. **Waking a parent never wakes its children.** Powering the controller doesn't mean anyone needs the chip.
3. **A parent can't sleep while a child is awake.**
4. **Device links add "depends on" arrows.** Each LU depends on the W-LUN (`device_link_add(…, DL_FLAG_PM_RUNTIME | DL_FLAG_RPM_ACTIVE)`, `ufshcd.c:5093`). So the W-LUN can't sleep while any LU is awake, and waking an LU wakes the W-LUN first.

On the board the LUs wait for 500 ms of quiet before they sleep (`MTK_RPM_AUTOSUSPEND_DELAY_MS`, `ufs-mediatek.h:367`, set as the SCSI host's `rpm_autosuspend_delay` at `ufs-mediatek.c:1769`), and the W-LUN waits 0 ms (`ufshcd_wl_probe()`, `ufshcd.c:10815`). So about half a second after the last disk access, the whole chip side is asleep.

TWO MEANINGS

Runtime suspend (one device sleeps while the phone is on) is not **system suspend** (the whole phone sleeps). This bug is entirely about runtime suspend.

And the one that caused confusion: `ufshcd_rpm_get_sync()` and `pm_runtime_get_sync()` are not two different mechanisms. The first is a one-line wrapper around the second:

```
static inline int ufshcd_rpm_get_sync(struct ufs_hba *hba)
{
        return pm_runtime_get_sync(&hba->ufs_device_wlun->sdev_gendev);
}
```

ufshcd-priv.h:344

Same function, different device: the W-LUN instead of `hba->dev`. Which device you name decides what wakes up.

ADDED 1 NOV · "SUSPEND" AND "RESUME" ARE LINUX WORDS

They come from Linux's power-management core (`runtime_suspend`, `runtime_resume`, `pm_runtime_get_sync()`). The specs name the same states differently:

WHAT	NAME IN THE SPEC	WHICH SPEC
The chip's power mode	Active, Sleep, PowerDown, DeepSleep, set with START STOP UNIT	JEDEC UFS (JESD220). SSU itself is a SCSI command; UFS defines the power-mode values it carries.
The link's sleep state	HIBERN8 is the M-PHY line state. UniPro calls it "Hibernate" and uses `DME_HIBERNATE_ENTER` and `EXIT`.	MIPI M-PHY and UniPro, not JEDEC
The hardware timer	Auto-Hibernate, the AHIT register	JEDEC UFSHCI (JESD223)

> The core translates between the two vocabularies with a table, `ufs_pm_lvl_states[]` (`ufshcd.c:243`). Each Linux power level is a pair: one chip power mode and one link state. The board's runtime level, 3, is {Sleep, hibern8} (line 247). So “Linux suspends the W-LUN” means START STOP UNIT to Sleep, plus the link into hibern8.
>
> The numbers carry straight through: `UFS_SLEEP_PWR_MODE` is 2, JEDEC's value for Sleep, and the core puts it in the command's power-condition field (`pwr_mode << 4`, `ufshcd.c:9516`).

3.4 What each object does when it sleeps and wakes (6.6)

This is where the real work happens. The code is short; read it. Comments starting with `←` are mine.

The W-LUN going to sleep: `__ufshcd_wl_suspend()`

```
        ret = ufshcd_vops_suspend(hba, pm_op, PRE_CHANGE);      // ← vendor hook
        ...
        if (req_dev_pwr_mode != hba->curr_dev_pwr_mode) {
                ...
                        ret = ufshcd_set_dev_pwr_mode(hba, req_dev_pwr_mode);
                                                // ← START STOP UNIT: chip to Sleep
        }
        ...
        ret = ufshcd_link_state_transition(hba, req_link_state, check_for_bkops);
                                                // ← DME_HIBERNATE_ENTER: link to hibern8
        ...
vops_suspend:
        ret = ufshcd_vops_suspend(hba, pm_op, POST_CHANGE);     // ← vendor hook
```

ufshcd.c:9746, lines 9830–9879 (trimmed)

How deep it sleeps is set by `rpm_lvl`. The core's default is level 3: chip in Sleep, link in hibern8 (`ufshcd.c:10676`). MediaTek switches to level 1 (chip stays Active, link in hibern8) only if the device tree has `mediatek,ufs-vcc-always-on`, which the MT6991 tree used here doesn't (`ufs-mediatek.c:2878`). Either way, **the link ends up in software hibern8**.

> **ADDED 1 NOV · WHAT MEDIATEK DOES IN ITS TWO SUSPEND HOOKS**
>
> The four steps above, with their lines:
>
> 1. `ufshcd.c:9830`: vendor `suspend`, `PRE_CHANGE`. MediaTek turns AHIT off and waits for the link to report UP.
> 2. `9840`: `ufshcd_set_dev_pwr_mode()` sends START STOP UNIT, and the chip goes to Sleep.
> 3. `9860`: `ufshcd_link_state_transition()` sends `DME_HIBERNATE_ENTER`, and the link is in software hibern8.
> 4. `9879`: vendor `suspend`, `POST_CHANGE`. MediaTek switches off the host's link hardware.
>
> The core does steps 2 and 3. Step 1 is MediaTek's addition; the stock core never sets AHIT to 0 on suspend. The core writes AHIT in only three places: at probe (`9012`), at the end of every W-LUN wake (`9999`), and from the sysfs `auto_hibern8` file.
>
> ### Step 1: why the timer goes off first
>
> From step 1 on, software drives the link by itself, and it needs the link in a known state:

- **START STOP UNIT travels over the link,** so the link must be UP to carry it. MediaTek's `ufs_mtk_auto_hibern8_disable()` (`ufs-mediatek.c:2997`) writes 0 to AHIT, waits up to 5 ms for the controller's state machine to go idle, then waits up to 100 ms for register 0x14C to say UP.
- **The core can't see auto-hibern8.** `hba->uic_link_state` changes only when software moves the link (`ufshcd.c:1905` in clock gating, `9608` in `ufshcd_link_state_transition()`). When the timer takes the link down, that variable still says ACTIVE. `__ufshcd_wl_suspend()` checks that the link is active before it starts (`9784`), but with the timer running that check would only be checking a variable.
- **The timer fires after 1 ms of quiet on the board.** If any gap between the steps lasted longer, the hardware would take the link down on its own, and software would then send its own hibernate-enter to a link that's already asleep.

Once the timer is off and the link is really UP, the hardware matches the core's variable, and nothing changes the link behind software's back. The comment in the function's failure branch shows it was written for this path: “trigger error handler and break suspend”.

Step 4: what `POST_CHANGE` switches off

By now the chip is asleep and the link is in software hibern8, so nothing will use the host's link hardware until the next wake. Everything below is on **the SoC's side**; the chip took care of its own insides when it received START STOP UNIT. In order (`ufs-mediatek.c:2627` to `2655`):

1. **UniPro goes into power-down.** UniPro is the protocol engine inside the host controller, and with the link asleep it has nothing to do. `ufs_mtk_link_set_lpm()` (`2521`) sets MediaTek's own UniPro attribute `VS_UNIPROPOWERDOWNCONTROL` (0xD0A8, `ufs-mediatek.h:128`) to 1.
2. **The M-PHY loses power.** These are the host's analog transmitters and receivers. In hibern8 they were already idle, with only a small wake-up detector on; now `phy_power_off()` removes their power (through `ufs_mtk_mphy_power_on()`, `607`).
3. **The SRAM is switched off.** This is a block of memory built into the UFS controller. Linux can't flip its power switch directly: `ufs_mtk_sram_pwr_ctrl()` makes an SMC, a call into MediaTek's secure firmware, which does it. (The macro is in mainline's `ufs-mediatek-sip.h`; the vendor tree makes the same call.) Neither MediaTek's code nor mainline says what the SRAM holds.
4. **The clock drops, at G5 only.** `_ufs_mtk_clk_scale(hba, false)` moves `ufs_sel` to the 273 MHz source and stops holding vcore at 0.65 V (`2655`).

A fifth action, holding the chip's reset line, runs only when the link is fully off, which doesn't happen on the board's runtime path. If any step fails, the hook marks the link as off and returns `-EAGAIN`, which forces a full reset instead of a normal wake later.

Waking undoes it in reverse. `ufs_mtk_resume()` (`ufs-mediatek.c:2669`) powers the SRAM back on, then the M-PHY, raises the clock if the link was at G5 (`2699`), then calls `ufs_mtk_link_set_hpm()` (`2467`). That function re-enables the controller (`ufshcd_hba_enable()`), takes UniPro out of power-down, sends `DME_HIBERNATE_EXIT`, and sets up the controller's request queues again. That last step is a sign the controller really lost its working state while asleep.

The W-LUN waking up: `__ufshcd_wl_resume()`

```
static int __ufshcd_wl_resume(struct ufs_hba *hba, enum ufs_pm_op pm_op)
{
        ...
        ret = ufshcd_vops_resume(hba, pm_op);                    // ← vendor hook
        ...
        if (ufshcd_is_link_hibern8(hba)) {
                ret = ufshcd_uic_hibern8_exit(hba);          // ← DME_HIBERNATE_EXIT
                if (!ret)
                        ufshcd_set_link_active(hba);
                ...
        } else if (ufshcd_is_link_off(hba)) {
                ret = ufshcd_reset_and_restore(hba);         // ← link was off: full restart
                ...
        }
        if (!ufshcd_is_ufs_dev_active(hba)) {
                ret = ufshcd_set_dev_pwr_mode(hba, UFS_ACTIVE_PWR_MODE);
                                              // ← START STOP UNIT: chip to Active
                ...
        }
        ...
        /* Enable Auto-Hibernate if configured */
        ufshcd_auto_hibern8_enable(hba);                         // ← AHIT timer back on
        ...
}
```

ufshcd.c:9927 (trimmed)

The controller waking up: `ufshcd_resume()`

```
static int ufshcd_resume(struct ufs_hba *hba)
{
        ...
        ufshcd_hba_vreg_set_hpm(hba);                // ← power supplies to full power
        ret = ufshcd_vreg_set_hpm(hba);
        ...
        /* Make sure clocks are enabled before accessing controller */
        ret = ufshcd_setup_clocks(hba, true);        // ← clocks on
        ...
        /* enable the host irq as host controller would be active soon */
        ufshcd_enable_irq(hba);                      // ← interrupt on

        goto out;
        ...
}
```

ufshcd.c:10163 (trimmed)

That's all the controller's wake-up does: power supplies, clocks, interrupt. Nothing about the link, nothing about the chip.

> **KEY POINT**
>
> In 6.6, the only wake-up that sends `DME_HIBERNATE_EXIT` is the W-LUN's. Waking the controller alone leaves the link asleep.

ADDED 1 NOV · THE CONTROLLER GOING TO SLEEP, AND MEDIATEK'S WRAPPER

The controller sleeps right after the W-LUN (rule 3). Its sleep function is the mirror image of `ufshcd_resume()` above: `ufshcd_suspend()` turns the interrupt off, turns the clocks off (`ufshcd.c:10136`), and then calls `ufshcd_vreg_set_lpm()` (`10147`) for the chip's supplies:

- chip not Active (it's in Sleep): VCC off;
- link in hibern8 or off: VCCQ into low-power mode. The voltage stays the same, but the regulator is told the load will be at most 1 mA (`UFS_VREG_LPM_LOAD_UA`), so the PMIC can switch it to a low-current mode. VCCQ stays on because the chip's controller has to keep its state.

MediaTek wraps this. Its `ufs_mtk_runtime_suspend()` (`ufs-mediatek.c:3615`) calls the core's version, puts more rails into low-power mode (`3626`), and then switches off the power domain the whole controller sits in (MTCMOS, `3629`; enabled by `mediatek,ufs-rtff-mtcmos`, `mt6991.dts:6764`). `ufs_mtk_runtime_resume()` (`3637`) does the reverse, then calls the core's wake.

So the controller's wake brings back the power domain, the supplies and the clocks. It doesn't bring back the chip's mode or the link: those belong to the W-LUN.

3.5 The sleep and wake order on the board

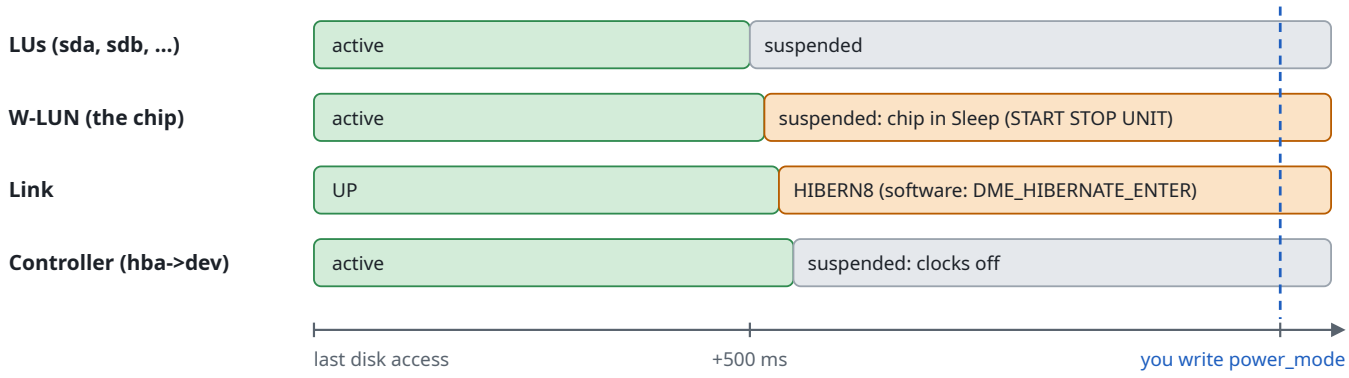


Figure 3. The sleep cascade. The LUs sleep after 500 ms of quiet; the W-LUN follows at once (chip to Sleep, link to software hibern8); then the controller turns its clocks off. By the time a command arrives, everything is asleep.

Going to sleep (500 ms without disk access): the LUs sleep → the W-LUN sleeps at once (chip to Sleep, link into software hibern8) → the controller sleeps (clocks off).

Waking with `ufshcd_rpm_get_sync()`, which names the W-LUN: rule 1 wakes the controller first (clocks on). Then the W-LUN's own wake-up runs: `DME_HIBERNATE_EXIT`, chip to Active, AH8 timer back on. The link is UP.

Waking with `pm_runtime_get_sync(hba->dev)`, which names the controller: clocks on. Then, by rule 2, nothing else. The link stays in hibern8 and the chip stays in Sleep.

UPDATE 1 NOV · TWO DETAILS IN THIS PICTURE

- **At +500 ms the link doesn't go straight from auto-hibern8 into software hibern8.** START STOP UNIT travels over the link, so the W-LUN's sleep first stops the timer and waits for the link to come back UP, then sends the command, and only then puts the link into software hibern8 (3.4).

- **“The controller sleeps (clocks off)” does more than clocks.** It also switches off the chip's VCC, puts VCCQ into low-power mode and, on the board, powers off the controller's whole power domain (3.4, 3.6).

3.6 The power ladder on the board ADDED 1 NOV

ADDED 1 NOV

Hibern8 is only the first rung. This is everything that switches off as the board stays quiet, in order:

QUIET FOR	WHO ACTS	WHAT GOES DOWN	THE CHIP
1 ms	The AHIT timer, in hardware	The link goes into hibern8: both ends' transmitters and receivers.	Active, all supplies on
10 ms	Clock gating (`ufshcd_gate_work()`, `ufshcd.c:1910`) and MediaTek's `setup_clocks` hook (`ufs-mediatek.c:893`)	The controller's clocks. Once MediaTek reads the link as in hibern8 in 0x14C, also the chip's reference clock and the host's M-PHY.	Active, all supplies on
~500 ms	`__ufshcd_wl_suspend()` (`ufshcd.c:9746`)	AHIT off, the chip into Sleep, the link into software hibern8. Then the host's UniPro, M-PHY and SRAM, and at G5 the clock (3.4).	Sleep, supplies on
right after	The controller's sleep: `ufshcd_suspend()` inside MediaTek's `ufs_mtk_runtime_suspend()`	Clocks off. Then the chip's VCC off and VCCQ into low-power mode. Then the controller's power domain (3.4).	Sleep, VCC off, VCCQ low-power

- **START STOP UNIT cuts no power.** It changes the chip's own mode: the chip finishes its work, stops using the flash, and becomes safe to lose VCC. Only the host switches the supplies, and it does that one rung later.
- **The 10 ms rung, in MediaTek's words:** “Gate ref-clk and poweroff mphy if link state is in OFF or Hibern8 by either Auto-Hibern8 or ufshcd_link_state_transition()”. The reference clock is the steady clock the host supplies to the chip's side of the link; it may stop while the link is in hibern8. The next command ungates the clocks, which powers the M-PHY and the reference clock back on, and then the doorbell wakes the link.
- **Waking climbs back down the ladder.** A normal read starts at the bottom and wakes everything (5.5). The old line woke only the last rung (5.4).

> KEY POINT
>
> The AHIT timer does rung 1. The W-LUN's sleep does START STOP UNIT plus its own software hibern8. The chip's supplies are cut one rung later, by the controller's sleep.

4 Why the original line was right once

Was this “decided in 6.6”? It wasn't. It changed in Linux 5.14, in 2021, and 6.6 inherited it. The ported code comes from Qualcomm's Linux 4.19, from before the change.

4.1 Up to Linux 5.13: one switch did everything

Before 5.14, waking `hba->dev` ran `ufshcd_runtime_resume()`, which called `ufshcd_resume(hba, UFS_RUNTIME_PM)`. In v5.13 that single function did all of it:

```
static int ufshcd_resume(struct ufs_hba *hba, enum ufs_pm_op pm_op)
{
        ...
        ret = ufshcd_setup_clocks(hba, true);                   // 8860 ← clocks on
        ...
        ufshcd_enable_irq(hba);                                 // 8865 ← interrupt on
        --------- 6.6 cuts the function here ---------
        ret = ufshcd_vops_resume(hba, pm_op);                   //      ← vendor hook
        ...
        if (ufshcd_is_link_hibern8(hba)) {
                ret = ufshcd_uic_hibern8_exit(hba);             // 8880 ← DME_HIBERNATE_EXIT
        ...
        if (!ufshcd_is_ufs_dev_active(hba)) {
                ret = ufshcd_set_dev_pwr_mode(hba, UFS_ACTIVE_PWR_MODE);
                                                                // 8905 ← chip to Active
        ...
        ufshcd_auto_hibern8_enable(hba);                        // 8928 ← AHIT timer back on
        ...
}
```

upstream v5.13, drivers/scsi/ufs/ufshcd.c:8846 (trimmed; numbers are line numbers)

So in 4.19, `pm_runtime_get_sync(hba->dev)` woke the controller, the link and the chip. Qualcomm's msm-4.19 debugfs function starts with exactly that line (`ufs-debugfs.c:1211` in msm-4.19, copied in file 04), and the ported code is a port of that function.

4.2 The deadlock that moved the job

The change is commit `b294ff3e3449`, “scsi: ufs: core: Enable power management for wlun”, by Asutosh Das and Can Guo (both at Qualcomm). It was committed on 10 May 2021 and first released in Linux 5.14. Its commit message explains what went wrong before it:

1. When the controller went to sleep, its own sleep function sent START STOP UNIT to the chip, to put the chip to sleep first.
2. But by rule 3, the controller only sleeps after all its children are asleep, and that includes the SCSI devices. Their request queues were already suspended.
3. A block-layer change, `52abca64fd94` (“scsi: block: Do not accept any requests while suspended”), made any request to a suspended queue wait until the queue resumes.
4. So the START STOP UNIT sat waiting for a queue that would never resume, and the controller's sleep waited for the command. Nothing could move. The commit message shows the stuck call chain:

`ufshcd_runtime_suspend → ufshcd_suspend → ufshcd_set_dev_pwr_mode → __scsi_execute → blk_get_request → blk_queue_enter → schedule`.

The fix gave the W-LUN its own runtime PM and moved the chip-and-link work into *its* sleep and wake functions. With the device links, the W-LUN sleeps after all the LUs and wakes after the controller, so its START STOP UNIT always has a working path.

4.3 The proof: the old function, cut in two

Put v5.13's `ufshcd_resume()` next to 6.6 and cut it right after `ufshcd_enable_irq()`. The top half is 6.6's `ufshcd_resume()`, the controller's wake-up. The bottom half is 6.6's `__ufshcd_wl_resume()`, the W-LUN's wake-up. The lines are nearly identical, as you can check against the excerpts in section 3.4.

STEP	V5.13: ALL IN `ufshcd_resume()`, RUN WHEN `hba->dev` WAKES	6.6: WHERE IT MOVED
Power supplies to full power	yes	`ufshcd_resume()`, the controller
Clocks on	line 8860	`ufshcd_resume()`, the controller
Interrupt on	line 8865	`ufshcd_resume()`, the controller
Vendor resume hook	yes	`__ufshcd_wl_resume()`, the W-LUN
`DME_HIBERNATE_EXIT`	line 8880	`__ufshcd_wl_resume()`, the W-LUN
Chip to Active (START STOP UNIT)	line 8905	`__ufshcd_wl_resume()`, the W-LUN
AHIT timer back on	line 8928	`__ufshcd_wl_resume()`, the W-LUN

LINUX 4.19 TO 5.13

`pm_runtime_get_sync(hba->dev)` wakes the controller, the link and the chip.

LINUX 5.14 AND LATER, INCLUDING 6.6

`pm_runtime_get_sync(hba->dev)` wakes the controller only.

PLAIN WORDS

In the old house, the main switch also turned on the hallway light. In 2021 the house was rewired and the hallway got its own switch. The ported instructions were written for the old house: “flip the main switch”. Flipping it left the hallway dark.

5 What happened on the board

5.1 The ported code path, start to finish

This is the ported code as of 31 October. Writing `echo 332211 > …/power_mode` runs:

1. **`ufsdbg_power_mode_write()`** reads the six digits, keeps the current rate (A or B), checks them with `ufsdbg_power_mode_validate()`, and compares them with the chip's maximum in `ufsdbg_cfg_pwr_param()`.
2. **`ufsdbg_config_pwr_mode()`** then:
   a. **Gets access to the hardware.** Originally with `pm_runtime_get_sync(hba->dev)`; now with `ufs_debugfs_get_user_access(hba)`.
   b. **`ufshcd_hold()`** keeps the controller's clocks on. Apart from runtime PM, the core also switches the controller's clocks off during short idle gaps. This is called **clock gating**, and on MediaTek it kicks in after at least 10 ms (`clk_gating.delay_ms = max(ah_ms, 10U)`). "Hold" blocks it while the code works.
   c. **`down_write(&hba->clk_scaling_lock)`** keeps the core's clock scaling, its error handler and its internal management commands (queries to the chip) from running at the same time. It's 6.6's replacement for msm-4.19's `hba->lock`.
   d. **`ufshcd_scsi_block_requests()`** tells the SCSI layer to stop sending new reads and writes.
   e. **`ufshcd_wait_for_doorbell_clr()`** waits until every command already sent has finished, that is, until the doorbell register (0x58) reads all zeros.
   f. **The ported clock lines** (section 6.3) and **`ufshcd_config_pwr_mode()`**, which runs MediaTek's `PRE_CHANGE` hook, then the PMC, then the `POST_CHANGE` hook.
   g. **Undo everything** in reverse order.

5.2 Error 1: the PMC timed out

This came from the first version, which, like Qualcomm's original (it calls `ufshcd_change_power_mode()` directly), sent the PMC without going through MediaTek's hooks. Whenever the board had been quiet for more than half a second, the W-LUN was asleep and the link was in software hibern8. The first line woke only the controller. The PMC was sent into a sleeping link, nothing answered, and after 500 ms the core gave up. It was first reported as "pmc command timed out"; the 6.6 core's wording is `pwr ctrl cmd … completion timeout`.

To be honest about this one: we never captured its full log. The sleeping link is the explanation that fits everything else, but it wasn't proven separately.

> **ADDED 1 NOV · WHY IT TIMED OUT, INSTEAD OF FAILING AT ONCE**
>
> - **The PMC doesn't try to wake anything.** The first line did the waking, and it reached only the controller (5.4). The PMC then went out into a link that was still asleep.
> - **A PMC is a negotiation, not a wake-up.** Writing `PA_PWRMODE` makes the host's UniPro send a power-change request across the wires and wait for the chip's confirmation (1.6). In software hibern8 the wires carry nothing until a hibernate-exit, so the request never crossed and the chip never answered.
> - **The message tells you which kind of failure it was.** The host accepted the command; otherwise `uic error` would have appeared at once (`ufshcd.c:4229`). What never came was the end of the exchange, so after 500 ms the result was `completion timeout` (`4237`).

- **What happened next.** The core marked the link broken and scheduled the error handler (`4270` to `4272`). The error handler wakes the W-LUN the right way, with `ufshcd_rpm_get_sync()` (`ufshcd.c:6381`), resets the controller, and brings the link back at maximum gear. So the next disk access worked, but the requested gear was gone.

5.3 Error 2: “exit h8 state fail”

Next the call was routed through `ufshcd_config_pwr_mode()`, the way MediaTek's own `clkscale_control` path does. That runs MediaTek's `PRE_CHANGE` hook *before* the PMC:

```
static int ufs_mtk_pwr_change_notify(struct ufs_hba *hba,
                                     enum ufs_notify_change_status stage,
                                     struct ufs_pa_layer_attr *dev_max_params,
                                     struct ufs_pa_layer_attr *dev_req_params)
{
        int ret = 0;
        static u32 reg;

        switch (stage) {
        case PRE_CHANGE:
                if (ufshcd_is_auto_hibern8_supported(hba)) {
                        reg = ufshcd_readl(hba, REG_AUTO_HIBERNATE_IDLE_TIMER);
                        ufs_mtk_auto_hibern8_disable(hba);    // ← return value ignored
                }
                ret = ufs_mtk_pre_pwr_change(hba, dev_max_params,
                                             dev_req_params);
                break;
        case POST_CHANGE:
                if (ufshcd_is_auto_hibern8_supported(hba))
                        ufshcd_writel(hba, reg, REG_AUTO_HIBERNATE_IDLE_TIMER);
                break;
        ...
```

ufs-mediatek.c:2023

And the function it calls:

```
int ufs_mtk_auto_hibern8_disable(struct ufs_hba *hba)
{
        unsigned long flags;
        int ret;

        /* disable auto-hibern8 */
        spin_lock_irqsave(hba->host->host_lock, flags);
        ufshcd_writel(hba, 0, REG_AUTO_HIBERNATE_IDLE_TIMER);
        spin_unlock_irqrestore(hba->host->host_lock, flags);

        /* wait host return to idle state when auto-hibern8 off */
        ufs_mtk_wait_idle_state(hba, 5);

        ret = ufs_mtk_wait_link_state(hba, VS_LINK_UP, 100);
        if (ret) {
                dev_warn(hba->dev, "exit h8 state fail, ret=%d\n", ret);

                spin_lock_irqsave(hba->host->host_lock, flags);
                hba->force_reset = true;
                hba->ufshcd_state = UFSHCD_STATE_EH_SCHEDULED_FATAL;
                schedule_work(&hba->eh_work);
                spin_unlock_irqrestore(hba->host->host_lock, flags);

                /* trigger error handler and break suspend */
                ret = -EBUSY;
        }

        return ret;
}
```

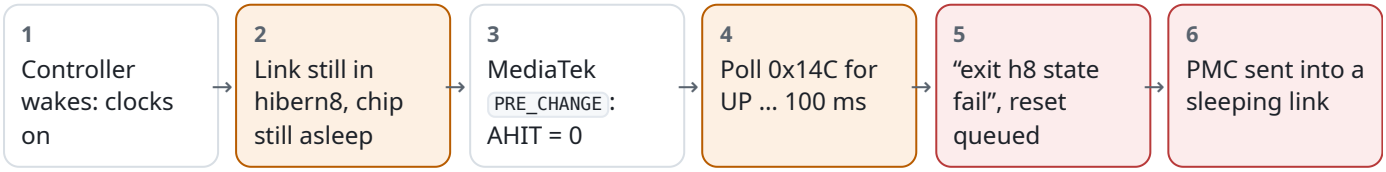
ufs-mediatek.c:2997

Step by step:

1. **Save AHIT and write 0**, which switches auto-hibern8 off. MediaTek's code expects that switching the timer off makes the controller bring the link out of *auto*-hibern8 by itself. That's why the next steps only wait and watch.
2. **Wait up to 5 ms** for any hibern8 entry or exit still in progress to finish.
3. **Poll register 0x14C for up to 100 ms**, waiting for the link state to read UP (2).
4. **It never did.** So it prints `exit h8 state fail, ret=-110` (−110 is `-ETIMEDOUT`), sets `force_reset`, marks the controller's state `UFSHCD_STATE_EH_SCHEDULED_FATAL` (3), schedules the error handler (which will reset the controller and the link) and returns `-EBUSY`.
5. **The caller ignores that return value** and goes on to `ufs_mtk_pre_pwr_change()` and then the PMC, on a link that's still asleep.

Why the link never came UP: it wasn't in auto-hibern8. The W-LUN's sleep had put it into *software* hibern8 with `DME_HIBERNATE_ENTER`. Switching the AHIT timer off can't undo that. Only `DME_HIBERNATE_EXIT` can, and the only code that sends it is the W-LUN's wake-up, which nothing had run.

Before the fix: `pm_runtime_get_sync(hba->dev)`



After the fix: `ufs_debugfs_get_user_access()` → `ufshcd_rpm_get_sync()`

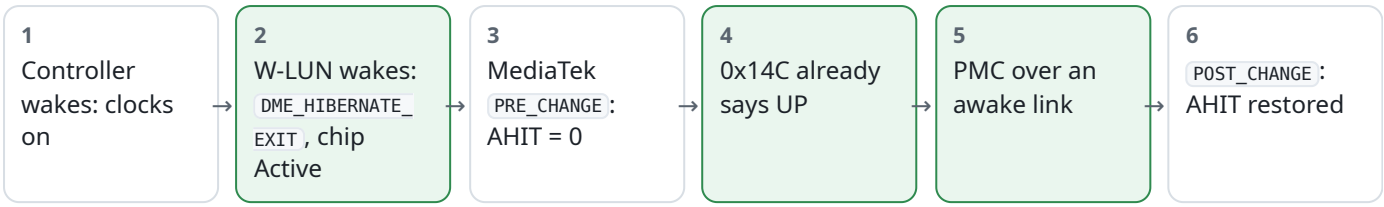


Figure 4. The same write, before and after. The only difference is step 2.

ADDED 1 NOV · WHY STOPPING THE TIMER COULDN'T HELP

MediaTek's hook assumes the link is either already UP or in auto-hibern8. In both cases, writing AHIT = 0 ends with the link UP (2.2). Software hibern8 is the case it doesn't cover:

- AHIT was already 0, because MediaTek's suspend hook had written it before START STOP UNIT (3.4). The write didn't even change the register.
- The hook waited 100 ms, the link never came up, and it printed `exit h8 state fail` and scheduled a reset (`ufs-mediatek.c:2997`).

The function's name says the same thing: `ufs_mtk_auto_hibern8_disable()` turns auto-hibern8 off. Taking the link out of software hibern8 isn't part of its job.

5.4 One bug, seen from two places

Errors 1 and 2 are the same problem. Without MediaTek's hook, the PMC itself ran into the sleeping link and timed out. With the hook, the hook noticed the sleeping link first and complained.

ADDED 1 NOV · THE BUG, ROW BY ROW

This uses the ladder from 3.6. When the command was typed, the board had been quiet for more than half a second, so every rung had already happened:

	BEFORE THE COMMAND	AFTER `pm_runtime_get_sync(hba->dev)`	AFTER `ufshcd_rpm_get_sync(hba)`
Controller	Power domain off, clocks off	On, clocks on	On, clocks on
Chip	Sleep, VCC off	VCC back on, **still in Sleep**	Active (START STOP UNIT)
Link	Software hibern8, AHIT = 0	**Still in software hibern8**	UP (`DME_HIBERNATE_EXIT`), AHIT back on

The old line names the controller, and by rule 2 waking a device never wakes its children. Nothing in the controller's wake sends a hibernate-exit or tells the chip to leave Sleep. Both are the W-LUN's job, and nobody woke the W-LUN. The new line names the W-LUN. By rule 1 that wakes the controller first, and then the W-LUN's own wake undoes its rung.

5.5 Why it could look random

Within half a second of any disk access, the LUs, and so the W-LUN, were still awake, the link was UP, and the old line was enough. If the old code ever worked, that's why.

> **UPDATE 1 NOV**
>
> “The link was UP” is too strong. Within half a second the W-LUN was awake, but after 1 ms of quiet the link was usually in *auto*-hibern8. That's still fine for the second version: MediaTek's hook stops the timer, and the controller brings the link back up by itself. For the first version, which skipped the hook, it's unclear: nothing shows whether a PMC on its own wakes an auto-hibernated link, and MediaTek's code doesn't rely on it.

> **ADDED 1 NOV · WHY NORMAL DISK USE NEVER HIT THIS BUG**
>
> Normal reads and writes come in from the bottom of the tree, and runtime PM wakes everything they depend on:
>
> 1. A read on `/dev/sda` goes to that LU's queue, and the block layer wakes the LU before sending the command.
> 2. The LU depends on the W-LUN through a device link with `DL_FLAG_PM_RUNTIME` (`ufshcd.c:5093`), so the W-LUN wakes first: hibernate-exit, chip to Active.
> 3. Before that, the W-LUN's parents bring the controller up.
>
> So every read wakes the controller, the chip and the link before its command goes out. The debugfs path came in from the top, at the controller, and waking a device never wakes anything below it.

5.6 The `#if 0` manual hibern8 exit

The ported code has a disabled block that called `ufshcd_uic_hibern8_exit()` directly. It aimed at the right symptom, but it would only have woken the link. The chip would still be in Sleep from START STOP UNIT, and the core's records would still say the W-LUN is suspended and the link is in hibern8, because `ufshcd_uic_hibern8_exit()` doesn't update `hba->uic_link_state`; its callers do (`ufshcd.c:4413`). The next real wake-up would then most likely send `DME_HIBERNATE_EXIT` to a link that's already awake, and fail. Waking the W-LUN does all three things, in the order the core expects.

Qualcomm's original also enters and exits hibern8 by hand, but for a different reason: around its clock switch, so that auto-hibern8 can't race it (`ufs-debugfs.c:1237` in msm-4.19).

6 The fix, and the other changes

6.1 The one line that mattered

Before (as ported from Qualcomm msm-4.19)

```
pm_runtime_get_sync(hba->dev);
ufshcd_hold(hba);
...
out:
...
ufshcd_release(hba);
pm_runtime_put_sync(hba->dev);
```

After (what runs now)

```
ret = ufs_debugfs_get_user_access(hba);
if (ret)
        return ret;
ufshcd_hold(hba);
...
out:
...
ufshcd_release(hba);
ufs_debugfs_put_user_access(hba);
```

The helper was already in the 6.6 file being edited:

```
static int ufs_debugfs_get_user_access(struct ufs_hba *hba)
__acquires(&hba->host_sem)
{
        down(&hba->host_sem);
        if (!ufshcd_is_user_access_allowed(hba)) {
                up(&hba->host_sem);
                return -EBUSY;
        }
        ufshcd_rpm_get_sync(hba);
        return 0;
}

static void ufs_debugfs_put_user_access(struct ufs_hba *hba)
__releases(&hba->host_sem)
{
        ufshcd_rpm_put_sync(hba);
        up(&hba->host_sem);
}
```

ufs-debugfs.c:68

- **`down(&hba->host_sem)`** takes the UFS core's “one big job at a time” semaphore. The error handler, system suspend, shutdown, the boot-time scan and the core's own sysfs writes all take it (`ufshcd.c` lines 6561, 10067, 10309, 9039, 1669), so the PMC can't collide with any of them.
- **`ufshcd_is_user_access_allowed()`** refuses during shutdown (it returns `!hba->shutting_down`).
- **`ufshcd_rpm_get_sync(hba)`** wakes the W-LUN, which by rule 1 wakes the controller first.

What happens now when you write `power_mode` after half a second of quiet:

1. The controller wakes: power supplies, clocks, interrupt.
2. The W-LUN wakes: `DME_HIBERNATE_EXIT` (link UP), START STOP UNIT (chip Active), AHIT timer back on.
3. MediaTek's `PRE_CHANGE` sets AHIT to 0; the link is already UP, so the check passes.
4. The PMC runs over an awake link. `POST_CHANGE` puts AHIT back.
5. The code releases everything, and half a second later the whole chip side sleeps again.

KEY POINT

MediaTek's own `clkscale_control` path never had this bug: it calls `ufshcd_rpm_get_sync()` before anything else (`ufs-mediatek.c:2322` and `2398`). The fix makes the path follow the same recipe.

ADDED 1 NOV · THE RULE THE FIX FOLLOWS, AND WHO ELSE FOLLOWS IT

The debugfs code never calls MediaTek's hook itself; `ufshcd_config_pwr_mode()` does. So the requirement falls on whoever calls `ufshcd_config_pwr_mode()`, and it comes down to two lines:

1. **`ufshcd_rpm_get_sync()` holds the W-LUN awake.** If the W-LUN is asleep, this wakes it, which takes the link out of software hibern8. If it's awake, this only stops it from sleeping until the hold is released. Either way, the link can't be in software hibern8 while the hold lasts.
2. **`ufshcd_hold()` keeps the clocks on.** It also covers the one other way into software hibern8, clock gating. That never happens on the board (2.2), and on boards where it does, ungating exits hibern8 before `ufshcd_hold()` returns (`ufshcd_ungate_work()`, `ufshcd.c:1751`).

Every caller in the kernel follows this: MediaTek's `clkscale_control` path (`ufs-mediatek.c:2322`), the error handler (`ufshcd.c:6381`), and probe and reset, which run their PMC right after link startup, on a link that has never slept. The old debugfs function was the only caller that skipped it.

Nothing enforces it. `ufshcd_config_pwr_mode()` doesn't check whether the W-LUN is awake; the rule lives only in how the callers are written. That's why the old code compiled and ran, and only failed in the hardware.

And `echo 4` is not part of the fix (6.4).

6.2 The other changes in file 06

Changed 1: the mode digits

The ported code had a “failsafe” block that forced FAST mode, using a variable called `pwr`. There is no `pwr` in the ported function, so it wouldn't compile, and forcing FAST would turn every SLOW (PWM) request into a FAST one. (The block looks like it came from an early version of file 05, which had both problems.) File 06 replaces it with:

```
        if (desired_pwr_mode->pwr_rx == FASTAUTO_MODE)
                desired_pwr_mode->pwr_rx = FAST_MODE;
        else if (desired_pwr_mode->pwr_rx == SLOWAUTO_MODE)
                desired_pwr_mode->pwr_rx = SLOW_MODE;
        desired_pwr_mode->pwr_tx = desired_pwr_mode->pwr_rx;
```

Why map the AUTO modes at all? MediaTek's `ufs_mtk_pre_pwr_change()` (`ufs-mediatek.c:1908`) asks a shared core helper, `ufshcd_get_pwr_dev_param()`, to agree the final settings. That helper only treats `FAST_MODE` as high speed:

```
        if (dev_max->pwr_rx == FAST_MODE)
                is_dev_sup_hs = true;
        ...
        if (!is_dev_sup_hs && is_pltfrm_max_hs) {
                pr_info("%s: device doesn't support HS\n", __func__);
                return -ENOTSUPP;
```

ufshcd-pltfrm.c (the same helper in a newer tree, renamed `ufshcd_negotiate_pwr_params()`; same logic)

A FASTAUTO request looks to it like “the device doesn't support HS”. It returns `-ENOTSUPP` and leaves the result all zeros. MediaTek only prints `failed to determine capabilities` and carries on, so the PMC would go out asking for gear 0 on 0 lanes. SLOW requests are fine, because MediaTek switches the helper into PWM mode for them.

The same helper also gives TX and RX the same gear (the lower of the two requested) and the same mode, caps lanes at 2, and sets rate B. So `342211` becomes G3 in both directions.

Changed 2: don't report success after a queued reset

```
        /* AH8 disable failed in MTK's PRE_CHANGE: host reset queued */
        if (!ret && hba->ufshcd_state != UFSHCD_STATE_OPERATIONAL)
                ret = -EIO;
```

MediaTek's hook can fail, queue a reset and still let the PMC report success. This turns that case into an error. The states (`ufshcd.h:529`):

VALUE	STATE	MEANING HERE
0	`UFSHCD_STATE_RESET`	Being reset
1	`UFSHCD_STATE_OPERATIONAL`	All good
2	`UFSHCD_STATE_EH_SCHEDULED_NON_FATAL`	Error handler queued for a minor error
3	`UFSHCD_STATE_EH_SCHEDULED_FATAL`	Reset queued. This is what MediaTek's failed check sets.
4	`UFSHCD_STATE_ERROR`	Recovery failed

The debug print, `dev_err(hba->dev, "state=%d\n", hba->ufshcd_state)`, reads the same field.

The validate function

Qualcomm's original `ufsdbg_power_mode_validate()` allows gears up to G4 and only FAST or FASTAUTO. It rejects G5 and every SLOW request with “power parameters are not valid”. File 06 has a version that allows G1 to G5 and all four modes, if you need it.

> **ADDED 1 NOV · IS THE WHOLE FUNCTION OK? A REVIEW**
>
> **Yes: the part that matters is right.** The question came up once the fio speeds started following the gear. It isn't luck:

- The ported function does the same steps, in the same order, as the core does before its own gear changes (`ufshcd_clock_scaling_prepare()`): wake, keep the clocks on, take `clk_scaling_lock`, block new commands, wait for the ones in flight, then the PMC.
- The wake is the core's own helper, the one its debugfs files already use (`ufs_debugfs_get_user_access()`, `ufs-debugfs.c:68`).
- The PMC goes through `ufshcd_config_pwr_mode()`, so MediaTek's hooks run. Those hooks send only link commands (DME), nothing that needs `clk_scaling_lock`, so holding that lock can't deadlock against them.

Three things still worth changing:

1. **Delete the ported clock block.** The reasons are in the 6.3 update below.
2. **Move the debug dumps.** The body of `ufshcd_dump_uic_state()` wasn't available. If it reads controller registers, two of its calls can run while the clocks are off: the “before hold” one, and the one in the access-failure path. MediaTek turns the clocks off 10 ms after the last command (`ufs-mediatek.c:1471`), but the W-LUN stays awake for about 500 ms. In that gap the wake does nothing and the clocks are off, and a register read there can hang the bus. Move the first call below `ufshcd_hold()`. Delete the failure-path one, which only runs during shutdown.
3. **Take the lock once per boot** (`echo 5`, then `echo 4`; see the 6.4 update).

> **WATCH OUT**
>
> Any UFS error resets the link to maximum gear and lanes without telling you. After each fio run, check `dmesg | grep "err_handler started"`. If that line appears, the run wasn't at the requested gear.

6.3 The ported clock-scaling lines

What they try to do: slow the controller's clock when the link goes down to G1, and speed it up for anything higher, by calling `ufshcd_scale_clks()` before or after the PMC.

What actually runs on MT6991:

1. `ufshcd_scale_clks()` first calls MediaTek's `clk_scale_notify` hook, which returns at once (`ufs-mediatek.c:3172`). The core's clock scaling is switched off on this board: the device-tree flag `mediatek,ufs-clkscale-only-by-user` (`mt6991.dts:6768`) makes MediaTek skip setting `UFSHCD_CAP_CLK_SCALING` (`ufs-mediatek.c:1763`).
2. Then the core's `ufshcd_set_clk_freq()` walks its list of clocks. `ufs_sel` is still on that list, with 273 and 499.2 MHz from the device tree; MediaTek removes only the two *source* clocks from the list (`ufs-mediatek.c:1240`). So the core calls `clk_set_rate(ufs_sel, …)`, unless it believes the clock is already at that rate.
3. MediaTek's own path makes the same switch differently (`_ufs_mtk_clk_scale()`, `ufs-mediatek.c:3039`). It raises vcore to at least 0.65 V *before* moving `ufs_sel` to the fast source with `clk_set_parent()`, and lowers vcore only *after* moving to the slow source. The core's path never touches vcore.

Is that safe? Slowing a clock down is always safe. Speeding it up to 499.2 MHz is safe only if vcore is already high enough. After `echo 4 > clkscale_control`, MediaTek has raised vcore, and the lock stops anyone lowering

it, so it's fine. The risky combination is `echo 1` or `echo 3` (MediaTek drops vcore to its minimum) followed by a G5 request through `power_mode`.

Two things not checked. Whether `clk_set_rate()` really moves this mux depends on MediaTek's clock driver, which wasn't read. And `ufshcd_scale_clks()` is `static` in the stock android15-6.6 `ufshcd.c` (line 1128); since the build links, the board's tree must make it visible, probably as part of the port.

CORRECTION

An earlier explanation said these lines were harmless because MediaTek's clock-scale hook does nothing. The hook does nothing, but the core still calls `clk_set_rate()` on `ufs_sel`. The conclusion stands (safe after `echo 4`; don't mix with 1 or 3), but the reason given was incomplete.

UPDATE 1 NOV · THE CLOCK BLOCK, LOOKED AT AGAIN

Two things missed earlier, and a recommendation.

MediaTek also moves the clock by itself, at G5

On the board MediaTek's suspend hook drops `ufs_sel` to 273 MHz whenever the link goes to sleep at G5, and its resume hook raises it again (`ufs-mediatek.c:2655`, `2699`). The ported function wakes the W-LUN first, so at G5 the clock is always back at 499.2 MHz before the PMC runs.

Where the risk really is

If nothing writes `clkscale_control`, the path stays consistent. We traced G5 → G1 → G5 → G4 with the W-LUN asleep before each write, and the clock is at 499.2 MHz whenever the gear is above G1. The risk needs a MediaTek scale-down (1 or 3) first. That moves `ufs_sel` to 273 MHz with `clk_set_parent()` and stops holding vcore at 0.65 V, without updating the core's record of the clock, `clki->curr_freq`. A G5 request through `power_mode` afterwards then goes one of two ways:

- If `curr_freq` still says 499.2 MHz, `ufshcd_scale_clks(hba, true)` sees nothing to do. The link runs G5 on the 273 MHz clock until the next sleep and wake.
- If `curr_freq` says 273 MHz (after a G1 visit), the core raises the clock with `clk_set_rate()` but never touches vcore. That's the case the text above warns about.

What either state actually does can't be told from the code: maybe it's only slower, maybe it causes link errors. MediaTek's own code never produces either.

Recommendation: delete the block

- **It does almost nothing on the board.** `ufs_sel` is the only clock in the table. For G2 to G5 the block either does nothing or undoes its own earlier G1 change.
- **Its one real action is a call MediaTek avoids.** For G1 it calls `clk_set_rate(ufs_sel, 273 MHz)`. MediaTek's comment above `ufs_mtk_clk_scale()` (`ufs-mediatek.c:3140`) explains their design: they switch the clock's source themselves first, so that the core's `clk_set_rate()` finds the rate already set and does nothing. The reason they give is to avoid “changing rate of pll clock that is shared between modules”. The port's G1 path is the one place where the clock framework picks the source by itself. It probably picks the 273 MHz one, and nothing seen so far suggests otherwise, but MediaTek's flow never exercises it.

- **Deleting it costs nothing.** The clock rule only works one way (1.3): G1 with the clock at 499.2 MHz is the normal state during every boot.
- **If you'd rather keep it,** check once: run `grep -iE "ufs|mainpll|univpll" /sys/kernel/debug/clk/clk_summary` before and after a G1 write. No PLL rate should change, and `ufs_sel`'s source should be `mainpll_d4_d2`.

> **PLAIN WORDS**
>
> “Don't gear changes need clock scaling?” Only if something can lower the clock. On the board three things can: this block itself (at G1), `clkscale_control` 1 or 3, and MediaTek's own sleep at G5, which its wake undoes. Delete the block and take the lock, and the clock never goes down, so it never needs to come back up. The block creates the very need it then fills.

6.4 `clkscale_control`, and why `echo 4` first

MediaTek's sysfs file takes six values. This is its own help text (`ufs-mediatek-sysfs.c:688`), with what each value does on the board:

VALUE	HELP TEXT	ON THE BOARD
0	free run	Refused: needs the core's clock scaling, which is off (“only by user”)
1	scale down	G4 and the 273 MHz clock (`min_gear = UFS_HS_G4`), vcore lowered
2	scale up	Max gear and 499.2 MHz. The value at boot.
3	scale down and forbid change	Like 1, then locked
4	scale up and forbid change	Like 2, then locked
5	allow change and free run	Removes the lock

What `echo 4` does, through `clkscale_control_store()` (`ufs-mediatek-sysfs.c:710`), `ufs_mtk_dynamic_clock_scaling()` (`ufs-mediatek.c:2364`) and `ufs_mtk_config_pwr_mode()` (`ufs-mediatek.c:2271`):

1. Wake the W-LUN (`ufshcd_rpm_get_sync()`), block new commands, hold the clocks, wait for the doorbell to clear.
2. Raise vcore, ask the CPUs for low wake-up latency (a PM QoS request), move `ufs_sel` to 499.2 MHz.
3. PMC to `max_gear`, then print `PMC Success! [RX, TX]: gear=…`.
4. Release everything, then set `clk_scale_forbid = true`.

Why it helps while you experiment:

- The clock and vcore are ready for G5.
- With the forbid flag set, every later call to `ufs_mtk_dynamic_clock_scaling()` returns at its first line. Its callers are writes to `clkscale_control` and the eye-monitor scan (`eyemon_scan()`). The store function accepts a

`powerhal_set:` prefix, which means Android's PowerHAL writes this file at runtime. Without the lock, one of those writes could run its own PMC and replace the setting.

WATCH OUT

`echo 2` or `echo 4` will not take you back to max gear after you've changed the gear through `power_mode`. `ufs_mtk_dynamic_clock_scaling()` remembers its last request in a `static` variable, `scale_mode`. If that was already “scale up”, it returns without a PMC; it doesn't know `power_mode` moved the gear. Write the max mode through `power_mode`, or reboot. `echo 5` removes the lock.

UPDATE 1 NOV · `ECHO 4` IS A LOCK, NOT PART OF THE FIX

Does the fix work without `echo 4`? It does. What fixed `exit h8 state fail` is the W-LUN wake at the top of the ported function, and `echo 4` can't be doing that job: MediaTek's handler wakes the W-LUN only for its own PMC and lets go at the end. With a 0 ms delay, the W-LUN is asleep again long before the next command. Test 3 (8.5) checks this on the board.

What the lock buys you

- **The requested gear can't be overwritten.** Any write of 1 to 4 runs MediaTek's own PMC: 1 and 3 set G4, 2 and 4 set the maximum gear. It copies `hba->pwr_info` (`ufs-mediatek.c:2280`), so it keeps the current lanes and mode and changes only the gear. With the lock set, every later write returns at the first check in `ufs_mtk_dynamic_clock_scaling()` (`2376`).
- **Two PMCs can't interleave.** MediaTek's path takes neither `host_sem` nor `clk_scaling_lock`, so a `clkscale_control` write while a `power_mode` write is running could mix its PMC steps with the other's.
- **No G5 on the slow clock.** The risky combination in the 6.3 update needs a 1 or 3, and the lock blocks both.

TRY IT: ONCE PER BOOT, BEFORE THE WRITES

```
C=/sys/bus/platform/devices/16810000.ufshci/clkscale/clkscale_control
echo 5 > $C
echo 4 > $C
```

The 5 comes first in case something already wrote 3 (scale down and lock). In that case a bare 4 hits the lock check and does nothing. On the board 5 never runs a PMC; it only clears the lock.

- **A free detector.** The `current:` line still changes when someone writes the file, even when the lock makes the write do nothing, because the store function records the value either way. If it later shows anything other than 4, something else on the system is writing this file.
- **About PowerHAL.** The text above says PowerHAL writes this file. That's an inference from the `powerhal_set:` prefix; it hasn't been confirmed on the build. The detector will tell.

What “only by user” means

The device-tree flag `mediatek,ufs-clkscale-only-by-user` (`mt6991.dts:6768`) means clock scaling happens only when *user space* asks for it, by writing `clkscale_control`. “User” is the kernel's word for anything

outside the kernel, PowerHAL included, not necessarily a person.

- MediaTek reads the flag into its own capability bit (`ufs-mediatek.c:826`). With that bit set, MediaTek doesn't turn on the core's clock scaling (`1763`), so the core never calls `ufshcd_devfreq_init()` (`ufshcd.c:8763`). There is no governor watching how busy the disk is.
- That's also why `clkscale_control` exists on the board: the `clkscale/` directory is created only if the core's scaling or this flag is on (`ufs-mediatek-sysfs.c:1061`). And it's why 0 (“free run”, hand control back to the governor) is refused (`ufs-mediatek-sysfs.c:736`): there is no governor to hand it to.
- MediaTek never sets G2 or G3 itself, and its only lower gear is G4. G1 appears only at boot (as PWM-G1) and in its debug code (1.3).

> **WATCH OUT**
>
> Don't write to `upthreshold` or `downdifferential`, the two files next to `clkscale_control`. They tune the governor that doesn't exist here, and both store functions call `mutex_lock(&hba->devfreq->lock)` without checking for NULL. `hba->devfreq` is NULL on the board, so writing a number to either file causes a kernel oops. Reading them is safe.

7 How we got here, and where we went wrong

7.1 The path, in order

1. **The question.** The starting point: the board, `clkscale_control`, the ported debugfs code and the two errors, and a search for a ready-made fix.
2. **The search.** We went through upstream Linux, MediaTek's vendor trees and the patch archives. No patch fixes this exact failure. Two are related, and neither fixes it: `7dec5cedf7b2` ("scsi: ufs: core: Serialize AHIT register access between sysfs and host driver") and `f29ec85ac4fd` ("scsi: ufs: host: mediatek: Return error directly on idle wait timeout"). We downloaded MediaTek's MT6991 vendor code (a public copy), the android15-6.6 core files and Qualcomm's msm-4.19 `ufs-debugfs.c`, so the three could be read side by side.
3. **One step at a time.** The first explanation was too compressed, so we slowed down and went through it a question at a time. What does disabling AH8 actually do? Must the link be up before a PMC? Why would the W-LUN's suspend have sent `DME_HIBERNATE_ENTER`? Why is there a separate step that powers UniPro down?
4. **Code to compare.** Files 01 to 05 were written into this folder, to open side by side in an editor.
5. **The flow.** The ported code's sequence: `pm_runtime_get_sync()`, hold, lock, block, doorbell, PMC, plus the ported clock lines. Its first step was the root cause: it woke `hba->dev`, not the W-LUN.
6. **The clock question.** Keep the ported clock-scaling lines, or copy MediaTek's? We settled on keeping them, setting `clkscale_control` to 4 first, and never mixing with 1 or 3.
7. **The ported code.** The code switched to `ufs_debugfs_get_user_access()` (from file 05) on 31 October. It was copied into file 06, with two more fixes marked.
8. **The test.** We rebuilt the kernel image, set `clkscale_control` to 4 and wrote `power_mode`: no `exit h8 state fail`. Then the lane tests passed too.

7.2 Corrections along the way

These statements changed during the investigation; knowing which helps in weighing the rest.

WHAT WAS SAID	WHAT'S ACTUALLY TRUE
Force every request to FAST mode.	Only the AUTO modes need mapping. Forcing FAST breaks SLOW (PWM) requests. (6.2)
First: the ported clock lines are dangerous. Then: they're harmless because MediaTek's hook does nothing.	They're safe after `echo 4`, but the core does call `clk_set_rate()` on `ufs_sel`. Don't mix them with 1 or 3. (6.3)
`down_write(&hba->lock)` is a problem.	That lock doesn't exist in 6.6. The 6.6 equivalent is `hba->clk_scaling_lock`.
Put the gear and lane function in MediaTek's vendor module (files 02, 03 and the patch).	The debugfs is in the core, which can't call vendor code or see `struct ufs_mtk_host`. Files 05 and 06 target the core. The vendor version is still there if you ever want a sysfs node; it was never compiled.

WHAT WAS SAID	WHAT'S ACTUALLY TRUE
An early file 05 used a variable named `pwr`.	The parameter is `desired_pwr_mode`. Fixed in 05 and 06.
The sequence was assumed to come straight from Qualcomm.	It came through a port for the MediaTek SoC, which took it from the Qualcomm SoC kernel.
“Nobody sends `DME_HIBERNATE_EXIT`.”	Imprecise. The W-LUN's wake-up sends it. It just never ran, because nothing woke the W-LUN.
“Since v5.14…”, which read as if 6.6 had decided it.	It changed in 5.14 (2021); 6.6 inherited it. (Part 4)
The link-state register was called `REG_UFS_PROBE`.	On MT6991 it's `REG_UFS_UFS_MMIO_OTSD_CTRL` (0x14C). `REG_UFS_PROBE` is the older chips' register. (2.3)

ADDED 1 NOV · MORE CORRECTIONS

WHAT WAS SAID	WHAT'S ACTUALLY TRUE
“`clkscale_control`, and why `echo 4` first” (6.4), and “with the fix, and `clkscale_control` at 4” (8.1). Easy to read as part of the fix.	`echo 4` isn't needed for the fix. It's a lock against other writers. Write 5, then 4. (6.4 update)
Earlier: after a controller-only wake, “the M-PHY stayed powered off”.	Not shown. The controller's wake runs MediaTek's `setup_clocks` hook, which calls `phy_power_on()`. Whether the PHY really came back depends on the PHY framework's reference count, which hasn't been traced. The explanation doesn't need it: software hibern8 keeps the wires silent until a hibernate-exit either way.
The risk analysis in 6.3.	It missed MediaTek's own clock moves at G5 on sleep and wake. (6.3 update)
6.4: “Android's PowerHAL writes this file at runtime.”	An inference from the `powerhal_set:` prefix, not confirmed on the build. (6.4 update)
5.5: “the link was UP”.	Within half a second the link was usually in auto-hibern8. (5.5 update)

KEY POINT

The bug wasn't in any one file. Each piece was right on its own terms: Qualcomm's line for Linux 4.19, the core's 2021 change, MediaTek's AH8 check for the case it was written for. The bug lived in the gap between them. Reading all three side by side is what found it.

8 How we know it's fixed

8.1 Results so far

- With the fix, and `clkscale_control` at 4: no `exit h8 state fail`.
- Lane changes from 2 to 1 pass.

Two more checks turn that into proof. If you've already run them, skip ahead.

> **ADDED 1 NOV · REPORTED SINCE**
>
> - The fio speeds follow the gear that was set. That's the physical check from 8.3, done with fio instead of `dd`.

8.2 Test 1: recreate the old failure condition

The old failure needed the W-LUN asleep. But `echo 4` wakes the W-LUN for a moment (its first scale-up after boot does), so a write right after it could pass even with the old code. Make sure the W-LUN is asleep before you write:

> **TRY IT**

```
W=/sys/bus/scsi/devices/0:0:0:49488/power/runtime_status   # the UFS Device W-LUN
cat $W                  # repeat until it says: suspended
cat $W; echo 112211 > /sys/kernel/debug/…/power_mode; dmesg | tail
```

> It passes if you see `suspended`, then no `exit h8 state fail`, and `state=1`. If `0:0:0:49488` isn't there, find it with `ls /sys/bus/scsi/devices | grep 49488`.

8.3 Test 2: did the gear really change?

“No error” isn't proof, for four reasons:

1. **The write never fails in the shell.** `ufsdbg_power_mode_write()` ends with `return cnt;` even when the PMC failed, so `echo` always looks successful. Only dmesg tells you.
2. **The success line in dmesg isn't from `power_mode`.** `PMC Success! [RX, TX]: gear=…` comes from MediaTek's `ufs_mtk_config_pwr_mode()`, that is, from `echo 4`. The core path prints nothing on success (`ufshcd_print_pwr_info()` uses `dev_dbg`, `ufshcd.c:697`).
3. **Asking for the current mode does nothing.** If the request equals the current settings, `ufshcd_change_power_mode()` returns 0 without sending a PMC (“power already configured”, `ufshcd.c:4566`). Right after `echo 4` the link is at max gear, so asking for max gear tests nothing.
4. **`cat power_mode` shows what the core believes.** It prints `hba->pwr_info`, which the core updates only after a successful PMC (`ufshcd.c:4638`). That's good evidence, but it's still software reporting on itself.

The physical check: the wire sets a speed limit. Whatever the software says, data can't move faster than the lanes carry it. The ceiling is bits per second × 0.8 (for 8b/10b) × number of lanes ÷ 8 bits per byte:

GEAR	BITS/S PER LANE	DATA PER LANE	CEILING, 2 LANES
HS-G1B	1,457.6 M	146 MB/s	~290 MB/s
HS-G3B	5,830.4 M	583 MB/s	~1.17 GB/s
HS-G4B	11,660.8 M	1,166 MB/s	~2.3 GB/s
HS-G5B	23,321.6 M	~2,870 MB/s (128b/130b)	~5.7 GB/s

These are ceilings before protocol overhead; real numbers land below them.

TRY IT

Go to G1 (`112211`), then:

```
echo 3 > /proc/sys/vm/drop_caches     # every run, or the read comes from RAM
dd if=/dev/block/sda of=/dev/null bs=1M count=1024
```

At G1, anything above about 290 MB/s means the link isn't at G1, because that's physically impossible. Then run the same at max gear and compare. Use G1 rather than G3: a simple `dd` may not fill a G3 link, so G3 and max gear could show similar numbers, while G1 is an unmistakable drop.

8.4 Reading the `state=` print

`state=1` (OPERATIONAL) means all good. `state=3` (EH_SCHEDULED_FATAL) means MediaTek's link check failed and a reset is queued: the old failure is back, so look at what woke (or didn't wake) the W-LUN.

8.5 Test 3: without `echo 4` ADDED 1 NOV

ADDED 1 NOV

This shows that the lock isn't part of the fix (6.4 update).

TRY IT

```
W=/sys/bus/scsi/devices/0:0:0:49488/power/runtime_status
# reboot, and don't touch clkscale_control
cat $W                  # repeat until it says: suspended
echo 112211 > /sys/kernel/debug/…/power_mode
dmesg | grep -E "exit h8|ufs_mtk_config_pwr_mode"
```

It passes if the `grep` prints nothing. No `exit h8` line means no failure. No `ufs_mtk_config_pwr_mode` line means MediaTek's own path never ran, so nothing else changed the gear; if one does appear, something wrote `clkscale_control`. Then check the gear the same way as in the lane test.

> **After every test run,** check `dmesg | grep "err_handler started"`. If it printed anything, an error reset put the link back to maximum gear during the run.

9 Open items

- Run tests 1 and 2, if you haven't.
- Make `ufsdbg_power_mode_write()` return the real error instead of `cnt` after a failed PMC, so `echo` fails visibly.
- To go back to max gear, write it through `power_mode` or reboot; `echo 2` or `echo 4` won't do it (6.4).
- If the failure ever comes back, print register 0x14C inside `ufs_mtk_auto_hibern8_disable()` at the moment it fails (needs a vendor-module rebuild). Bits 31 to 24 reading 3 would confirm a link stuck in hibern8.
- The board's vendor tree and the public copy used here have the same vendor code, but check line numbers before relying on them.
- `ufsdbg_cfg_pwr_param()` writes `hba->max_pwr_info.info.pwr_rx = FAST_MODE`. That's core state the error handler uses to restore the link. It's harmless today, because the chip already reports FAST there; don't write anything else into it.

> **ADDED 1 NOV · NEW OPEN ITEMS**
>
> - Decide on the ported clock block: delete it, or keep it after the `clk_summary` check (6.3 update).
> - Move or delete the two `ufshcd_dump_uic_state()` calls that can run with the clocks off (6.2 review).
> - Make `echo 5` then `echo 4` a boot step, and watch the `current:` line to see whether anything else writes `clkscale_control` (6.4 update).
> - Run test 3 (8.5).
> - Never write `upthreshold` or `downdifferential` on this board (6.4 update).
> - Optional: trace the PHY framework's reference count, to settle whether a controller-only wake powers the M-PHY back on (7.2).

10 Glossary

AH8 (auto-hibern8)

The controller puts the link into hibern8 by itself after an idle time, and wakes it by itself when new commands arrive.

AHIT

The controller register (offset 0x18) that holds the AH8 idle time. 0 means off.

`clk_scale_forbid`

MediaTek flag set by `clkscale_control` values 3 and 4. It blocks MediaTek's own clock and gear changes until value 5 clears it.

Clock gating

Switching the controller's clocks off during short idle gaps (at least 10 ms on MediaTek). Separate from runtime PM. `ufshcd_hold()` blocks it.

DME

The part of UniPro that takes management commands from the driver: `DME_GET`, `DME_SET`, `DME_HIBERNATE_ENTER`, `DME_HIBERNATE_EXIT`.

Doorbell

Controller register (offset 0x58) with one bit per command slot. The driver sets a bit to start a command; the controller clears it when the command is done.

EH (error handler)

The core's recovery work. It resets the controller and restarts the link.

Gear

Bits per second per lane. HS-G1 to G5 are the fast family; the PWM gears are the slow one.

GKI

Android's Generic Kernel Image, the common kernel that holds the UFS core.

`hba`

“Host bus adapter”: the core's `struct ufs_hba`, describing one controller. `hba->dev` is the controller's `struct device`.

Hibern8

The link's deepest sleep that keeps its settings.

`host_sem`

The core's semaphore that keeps user requests, error handling, system suspend and shutdown from overlapping.

Lane

One pair of wires carrying data in one direction.

LU

Logical unit: a readable, writable slice of the chip (`sda`, `sdb`, ...).

M-PHY

The electrical layer: the wires, their voltages, and recovering bits and timing from them.

PMC

Power mode change: switching gear, lanes or mode.

Rate A, rate B

The two base frequencies for HS gears. MediaTek uses B.

Runtime PM

Linux's per-device sleep and wake while the system is running.

SSU (START STOP UNIT)

The SCSI command that sets the chip's power mode: Active, Sleep, PowerDown.

UIC command

A command for the link layer, written to controller register 0x90.

UniPro

The link layer: frames, acknowledgements, re-sending, speed and power state.

UPIU

The packet format that carries SCSI commands over UFS.

vcore

The SoC's core supply voltage. The 499.2 MHz `ufs_sel` source needs at least 0.65 V.

vops

The table of vendor hooks that the core calls.

W-LUN

Well-known logical unit. The UFS Device W-LUN (`0xD0`, shown by Linux as `0:0:0:49488`) stands for the whole chip.

> ADDED 1 NOV · MORE TERMS
>
> ### Device link
>
> A “depends on” arrow between two Linux devices. Each LU depends on the W-LUN, so waking an LU wakes the W-LUN first.
>
> ### devfreq
>
> Linux's framework for changing a device's frequency according to load. The core's UFS clock scaling uses it. It's never started on the board (“only by user”).
>
> ### MTCMOS
>
> MediaTek's name for a switchable power domain inside the SoC. On the board MediaTek switches off the UFS controller's domain after the controller sleeps.
>
> ### PACP
>
> The control frames that UniPro's lowest layer sends over the link. A PMC is a `PACP_PWR_req` from the host and a `PACP_PWR_cnf` back from the chip.

PMIC

Power-management chip: the separate chip (MT6363 here) whose regulators make the board's supply rails, including the UFS chip's VCC and VCCQ.

Power level (`rpm_lvl`)

The core's pairing of a chip power mode with a link state, from `ufs_pm_lvl_states[]`. The board's runtime level, 3, is {Sleep, hibern8}.

Reference clock

The steady clock the host supplies to the chip's side of the link. MediaTek stops it once the link is in hibern8.

`setup_clocks`

The vendor hook the core calls whenever it turns the controller's clocks on or off: in clock gating, and in the controller's own sleep and wake.

SMC

Secure monitor call: an instruction that asks the secure firmware to do something Linux isn't allowed to do directly. MediaTek uses one to power the UFS SRAM on and off.

SRAM

On-chip memory. The UFS controller has a block of it, which MediaTek switches off when the chip side sleeps. The code doesn't say what it holds.

Symbol clock

A clock running at the link's own rate: `unipro_tx_sym` for sending, and one receive clock per lane (`unipro_rx_sym0` , `unipro_rx_sym1`), recovered from the arriving bits.

`uic_link_state`

The core's record of the link's state. Only software moves update it; auto-hibern8 doesn't.

UniPro power-down

MediaTek's own UniPro attribute, `VS_UNIPROPOWERDOWNCONTROL` (0xD0A8). It's set to 1 when the chip side sleeps, so the host's UniPro engine stops.

VCC

The supply rail for the chip's flash memory. The host switches it off once the chip is in Sleep.

VCCQ

The 1.2 V supply rail for the chip's controller and its end of the link. On the board's runtime sleep it goes into low-power mode, never off.

11 Files in this folder

`~/Personal_Study/kernel/mt6991_ufs_gear_lane/`

FILE	WHAT IT IS
`01_mtk_reference.c`	MediaTek's own `ufs_mtk_config_pwr_mode()`, with its steps marked [1] to [8]. The correct recipe, for reference.
`02_ufs_mtk_config_gear_lane.c`	A gear and lane function for the vendor module. Alternative path; not compiled.
`03_gear_lane_sysfs.c`	A sysfs file that calls 02.
`0001-ufs-mediatek-add-gear_lane-sysfs.patch`	02 and 03 as a patch for the vendor tree. Applies to the public copy; not compiled.
`vendor_6.6/`, `patched_6.6/`	The three MediaTek files before and after that patch.
`04_qcom_msm-4.19_ufsdbg_config_pwr_mode.c`	Qualcomm's original, which was ported to the MediaTek SoC.
`05_core_ufs-debugfs_config_pwr_mode.c`	A core version built on `ufs_debugfs_get_user_access()`.
`06_ported_ufsdbg_config_pwr_mode_fixed.c`	The ported function, with the fixes marked `CHANGED`.
`UFS_gear_lane_fix_explained.html` and `.pdf`	This document.

ADDED 1 NOV · NEW FILES

FILE	WHAT IT IS
`sources/`	The source copies this document cites, so every `file:line` can be looked up: `android15-6.6/` (the core: `ufshcd.c`, `ufshcd.h`, `ufshcd-priv.h`, `ufs-debugfs.c`), `mtk_vendor_6.6/` (the MediaTek files and `mt6991.dts`), `upstream_v5.13_v5.14/` (Part 4) and `qcom_msm-4.19/` (Qualcomm's `ufs-debugfs.c`). Two catches: `ufshcd.c:8846` in Part 4 is in the v5.13 file, and `ufs-debugfs.c:1211` and `1237` are Qualcomm's.
`sources/README.md`	Which citation points to which file, plus the PMC call tree from 1.6.
`UFS_gear_lane_fix_explained_v1.html` and `_v1.pdf`	The first version of this document, from 31 October, unchanged.

12 References

Commits

- `b294ff3e3449`: “scsi: ufs: core: Enable power management for wlun”. Asutosh Das and Can Guo; committed 10 May 2021; in Linux 5.14. The change at the heart of Part 4.
- `52abca64fd94`: “scsi: block: Do not accept any requests while suspended”. The block-layer change that turned the old design into a deadlock.
- `7dec5cedf7b2`: “scsi: ufs: core: Serialize AHIT register access between sysfs and host driver”. Related; doesn't fix this.
- `f29ec85ac4fd`: “scsi: ufs: host: mediatek: Return error directly on idle wait timeout”. Related; doesn't fix this.

Source copies read

- MediaTek vendor code: a public MT6991 kernel release, `kernel/kernel_device_modules-6.6/drivers/ufs/` (`ufs-mediatek.c`, `ufs-mediatek.h`, `ufs-mediatek-sysfs.c`) and `mt6991.dts`.
- Android common kernel `android15-6.6`: `drivers/ufs/core/ufshcd.c`, `ufs-debugfs.c`, `ufshcd-priv.h`, `include/ufs/ufshcd.h`.
- Upstream Linux v5.13 and v5.14: `drivers/scsi/ufs/ufshcd.c`.
- Qualcomm msm-4.19 (LineageOS sm8250 kernel): `drivers/scsi/ufs/ufs-debugfs.c`.
- The ported code in `kernel-6.6/drivers/ufs/core/ufs-debugfs.c`, 31 October 2026.

> **ADDED 1 NOV**
>
> - The copies above are now in `sources/` (Part 11). The MediaTek folder also has `ufs-mediatek-dbg.c`, `ufs-mediatek-perf.c`, `ufs-mediatek-priv.h` and `phy-mtk-ufs.c`.
> - Mainline Linux (the 6.18 tree in `~/Personal_Study/kernel/linux`): `drivers/ufs/host/ufs-mediatek-sip.h`, where `ufs_mtk_sram_pwr_ctrl()` becomes an SMC, and `include/ufs/ufs.h` for `UFS_VREG_LPM_LOAD_UA` (1000 µA).

Written with Claude Code, 31 October 2026, from a debugging session with Shrihari. All line numbers refer to the source copies listed above. Updated 1 November 2026 with answers to follow-up questions; nothing from the first version was removed.